



TECHNICAL UNIVERSITY

OF CLUJ-NAPOCA, ROMANIA

FACULTY OF AUTOMATION AND COMPUTER SCIENCE

Real-Time Light Dosimetry for Heritage Digital Twins via Texture Space Ray Tracing

LICENSE THESIS

Graduate: **Sebastian Duică**

Supervisor: **Assoc. Prof. Victor Ioan Băcu**

2026



TECHNICAL UNIVERSITY

OF CLUJ-NAPOCA, ROMANIA

FACULTY OF AUTOMATION AND COMPUTER SCIENCE

DEAN,
Prof. dr. ing. Vlad MUREȘAN

HEAD OF DEPARTMENT,
Prof. dr. ing. Rodica POTOLEA

Graduate: **Sebastian Duică**

Real-Time Light Dosimetry for Heritage Digital Twins via Texture Space Ray Tracing

1. **Project proposal:** Development of ChronoLux, an interactive Digital Twin for real-time light dosimetry using Texture Space Ray Tracing in DirectX 12.
2. **Project contents:** Introduction, Project Objectives, Bibliographic Research, Analysis and Theoretical Foundation, Detailed Design and Implementation, Testing and Validation, User's Manual, Conclusions, Bibliography.
3. **Place of documentation:** Technical University of Cluj-Napoca, Computer Science Department
4. **Consultants:**
5. **Date of issue of the proposal:** October 1, 2025
6. **Date of delivery:** July 10, 2026

Graduate: _____

Supervisor: _____

**TECHNICAL UNIVERSITY**

OF CLUJ-NAPOCA, ROMANIA

FACULTY OF AUTOMATION AND COMPUTER SCIENCE

**Declarație pe proprie răspundere privind
autenticitatea lucrării de licență**

Subsemnatul(a) _____, _____ legitimat(ă) cu

_____ seria _____ nr. _____
(conform și copiei anexate prezentei declarații, copie sem-
nată și certificată "conform cu originalul"), autorul lucrării

elaborată în vederea susținerii examenului de finalizare a studiilor de licență la Facultatea de Automatică și Calculatoare, Programul de studiu _____ din cadrul Universității Tehnice din Cluj-Napoca, sesiunea _____ a anului universitar _____, declar pe proprie răspundere, că această lucrare este rezultatul propriei activități intelectuale, pe baza cercetărilor mele și pe baza informațiilor obținute din surse care au fost citate, în textul lucrării și în bibliografie.

Declar, că această lucrare nu conține porțiuni plagiate, iar sursele bibliografice au fost folosite cu respectarea legislației române și a convențiilor internaționale privind drepturile de autor.

Declar, de asemenea, că această lucrare nu a mai fost prezentată în fața unei alte comisii de examen de licență.

În cazul constatării ulterioare a unor declarații false, voi suporta sancțiunile administra- tive, respectiv, *anularea examenului de licență*. Sunt de acord ca, pe tot parcursul vieții, în cazul în care este necesar și se va dori verificarea autenticității lucrării mele să fiu identificat și verificat în baza datelor declarate de mine și conform copiei documentului de identitate menționat mai sus.

Data

Nume, Prenume

Semnătura

General Advice

READ FIRST (this page should be removed from the final version):

1. If you print the thesis on paper, the three preceding pages (title page, summary page and declaration) should be printed on separate pages (one-side) and should be included in the printed paper. The summary page (the second one) must be signed by the graduate and the supervisor. The date on the declaration should be the date when the thesis is handed to the commissions' secretary.
2. On the title page, you should include the correct scientific title of your supervisor (use the Department web pages to find that).
3. Every chapter begins on a new page.
4. Do not change page margins.
5. Obey the other instructions provided in every chapter.
6. This template obeys all the formatting rules.
7. We included the `hyperref` package to generate navigation links. To produce a version for printing on paper uncomment the line containing `%\hypersetup{hidelinks}` located in the beginning of the main file `thesis_eng.tex`.

Contents

Chapter 1	Introduction	1
1.1	Context and Motivation	1
1.2	Limitations of Existing Solutions	1
1.3	The Proposed Solution: ChronoLux	2
1.4	Thesis Structure	2
Chapter 2	Project Objectives	4
2.1	Technical and Architectural Objectives	4
2.2	Metrological and Validation Objectives	4
Chapter 3	Bibliographic Research	6
3.1	Preventive Conservation and Empirical Dosimetry	6
3.2	The Metrology of Global Illumination	7
3.2.1	Rasterization vs. Ray Tracing in Scientific Simulation	7
3.2.2	The Rendering Equation and Monte Carlo Integration	7
3.2.3	Physical Material Simulation and Energy Conservation	8
3.2.4	Limitations of Offline Rendering	9
3.3	Hardware Acceleration and Heritage Digital Twins	9
3.3.1	Bounding Volume Hierarchies and Ray Traversal	9
3.3.2	DirectX Raytracing (DXR)	9
3.3.3	Variance Reduction and Texture Space Rendering	10
3.3.4	Solar Telemetry and Ephemeris Models	10
3.3.5	The Perez All-Weather Sky Model	11
Chapter 4	Analysis and Theoretical Foundation	12
4.1	High-Level Theoretical Pipeline	12
4.1.1	The Initialization Domain	12
4.1.2	The Radiometric Integration Domain	13
4.1.3	The Temporal Accumulation Domain	13
4.2	The Texture Space Ray Tracing Paradigm	13
4.2.1	The Mathematical Flaws of Screen Space Metrology	13
4.2.2	Texture Space Mapping	14
4.3	Dosimetric Path Tracing Algorithm	14
4.3.1	Stochastic Hemisphere Sampling and Energy Attenuation	14
4.3.2	The Physics of Refractive Transmission	15
4.3.3	Next Event Estimation (NEE)	15
4.3.4	Miss Evaluation and Perez Integration	16
4.4	The Parallel Determinism Problem	16
4.5	Mathematical Dose Accumulation	17
4.5.1	Temporal Integration Logic	17
4.5.2	Floating-Point Precision and Heatmap Generation	17
4.6	Conclusion	17

Chapter 5	Detailed Design and Implementation	19
5.1	Software Architecture Overview	19
5.2	Core Simulation Engine (C# CPU Orchestration)	19
5.2.1	The Chronological Engine and Time Stepping	21
5.2.2	Mathematical Astrometry and Solar Ephemeris	21
5.3	Scene Management and Hardware Interfaces	22
5.3.1	Asynchronous Runtime Mesh Parsing	22
5.3.2	Dynamic Material Serialization and Hierarchy Keying	23
5.3.3	Spatial Normalization	23
5.3.4	Deterministic BVH Building	24
5.4	Hardware-Accelerated Path Tracing (HLSL / DXR)	24
5.4.1	Memory Management and DXR Acceleration Structures	24
5.4.2	Texture Space Dispatch Paradigm	25
5.4.3	Pseudo-random Sequence Generation and Hashing	26
5.4.4	Direct Sunlight and Stochastic Multi-Sample Paths	27
5.5	Advanced Mathematical Models and Optimization	27
5.5.1	Perez Sky Model Atmospheric Coefficients	28
5.5.2	Hardware Traversal Mechanics (DXR and Face Culling)	28
5.6	Asynchronous Data Telemetry and UI Integration	28
5.6.1	Automated Sensor Placement	28
5.6.2	Metrological Analytics Pipeline	29
5.6.3	User Interface Toolkit and JSON Serialization Limitations	30
5.7	Deep Engine Optimizations and Memory Integrity	30
5.7.1	Ray-Origin Epsilon Offsetting and Floating Point Precision Limitations	31
5.7.2	Helmholtz Reciprocity and Energy Conservation Enforcement	31
5.7.3	Thread Synchronization and CPU-GPU Data Dependencies	31
5.7.4	VRAM Memory Integrity and Cleaning Protocol	32
5.7.5	Data Export Payload Calculation and Mathematical Variance	32
Chapter 6	Testing and Validation	33
6.1	Introduction to the Metrological Validation Framework	33
6.2	Experimental Setup and Simulation Parameters	33
6.2.1	The Virtual Environment and Target Artifact	33
6.2.2	Meteorological and Astrometric Configuration	33
6.2.3	Atmospheric Modeling via the Perez Model	34
6.2.4	Hardware and Execution Profiling	34
6.3	Data Telemetry and the Simulation Output	34
6.3.1	Downsampled Dosemap and GPU Readback	34
6.3.2	The Telemetry CSV Format	35
6.4	Overview of Conservation Scenarios	35
6.5	Scenario A: The Baseline Exposure	36
6.5.1	Configuration and Untreated Conditions	36
6.5.2	Empirical Analysis of the Catastrophic Dose	36
6.5.3	Diurnal Fluctuation and Hourly Stress	36
6.6	Scenario B: Material Intervention	38
6.6.1	Simulating Physical Filtration	38
6.6.2	The Physics of Ray Attenuation	38

6.6.3	Results of the Material Intervention	38
6.7	Scenario C: Positional Optimization and Geometric Occlusion	39
6.7.1	The Role of Architectural Shadows	39
6.7.2	Phenomenological Analysis: Window vs Corner Geometry	39
6.8	Scenario D: Combined Optimal Conservation	40
6.8.1	Synergistic Interventions	40
6.9	Scientific Validation: Variance and Convergence	40
6.9.1	The Necessity of Mathematical Integrity	40
6.9.2	Spatial Dose Variance Tracking	41
6.9.3	Monte Carlo Convergence Error (SEM)	41
6.9.4	Asynchronous Telemetry Pipeline Integrity	42
6.10	Conclusion of the Empirical Validation	43
Chapter 7	User's Manual	44
7.1	Introduction	44
7.2	Installation and System Requirements	44
7.2.1	Hardware Prerequisites	44
7.2.2	Software Installation and Artifact Import	44
7.3	The Project Launcher Dashboard	45
7.4	Viewport Navigation and Interaction	45
7.5	The Laboratory HUD and Telemetry	46
7.6	Material Configuration and Catalog	46
7.7	Simulation Configuration and Execution	48
7.8	Heatmap Visualization and Data Export	48
7.8.1	Viewport Heatmap	48
7.8.2	Headless Telemetry Export	48
Chapter 8	Conclusions	49
Bibliography		51
Chapter 9	Appendix	53
9.1	Core HLSL Path Tracing Implementation	53
9.2	C# Orchestrator Implementations	55
9.3	Mathematical Radiometry and Ephemeris Models	57

Chapter 1. Introduction

1.1. Context and Motivation

The concept of cultural heritage conservation is one that is highly complex, being highly scientific and involving the balancing act between access to the public and the protection of artifacts. In museums, it is common practice to face a situation where there is a need to display old paintings, which could be as old as 16th century oil paintings, delicate pieces of cloth and even water colors without damaging them. The most ubiquitous source of degradation of such artifacts, however, is photochemical degradation caused by the presence of light.

Light Dosimetry is the metrological way of quantifying the radiative energy, measured in Lux Hours, acting upon an artifact within a specific period. Given that this damage is strictly cumulative and irreversible, it becomes imperative that preventive conservation measures be taken. For this purpose, guidelines set forth by international bodies like the International Commission on Illumination (CIE) are used to govern preventive conservation of artifacts. For highly susceptible items, the recommendation made by CIE is that the maximum limit of energy to which the artifact may be exposed should not exceed 50,000 Lux Hours annually.

1.2. Limitations of Existing Solutions

Up to now, the management of the stress induced by the light had been relying on two separate paradigms. However, both have their respective limitations when applied to today's workflows in the conservation of valuable heritage.

According to the first approach, physical light measurement can be carried out by the means of IoT sensor networks and photometers. While these devices indeed provide the essential ground truth data, they cannot help predict a threat before its occurrence; therefore, there is always going to be some delay, which makes these sensors unable to provide any proactive measure prior to a new exhibition. Besides, while measuring irradiance at a single point in space, it would take a huge amount of sensors deployed at high density in order to fully cover all possible angles of light incidence in a large object. Not only it would be expensive and unattractive, but mounting these devices on fragile pieces would simply cause irreversible damage. Finally, as every other device with a physical sensor inside, it suffers from calibration error, which needs to be compensated for by regular maintenance.

According to the second one, the issue of predicting is approached by using offline rendering engines like Radiance or Precomputed Radiance Transfer (PRT) lightmaps, which calculate the lighting distribution mathematically and in an extremely reliable way. Yet their inability to exploit hardware acceleration capabilities and being completely CPU-bound renders them impractical for real-time applications due to their poor

performance, which results in hours' worth of calculation just for a single, one-year exposure simulation done with hourly precision (which is required in order to achieve proper global illumination).

That is precisely why this offline approach does not allow any interactive exploratory process to be carried out, meaning that conservators can neither explore the possibility of introducing some interventions into the environment—like altering transmittance of window glass, changing room albedo or moving objects around—nor try some optimal combination thereof.

1.3. The Proposed Solution: ChronoLux

For mitigating the issue of the interactivity and fidelity gap, this research will introduce ChronoLux - an interactive Digital Twin platform designed specifically for near-real-time dosimetry. Making use of modern GPU acceleration capabilities through DirectX 12 Raytracing (DXR), ChronoLux will transform passive artifact three-dimensional representations into interactive, predictive, and accurate simulations.

As opposed to camera-oriented screen-space rendering through camera-to-image ray-casting performed by current generation interactive graphics platforms, ChronoLux introduces a novel architectural approach - the Texture Space Ray Tracing. Regular renderers conduct frustum culling to exclude rays that are cast beyond the view range of the camera. Such an approach is insufficient when simulating light interaction for scientific dosimetry since measurements need to be conducted persistently throughout the entire geometrical surface of an object. Instead, using the artifact's UV coordinates, ChronoLux will cast rays from those points thus baking any cumulative exposure permanently into the object's surface.

In order to guarantee accuracy and metrological precision, the simulation will run under the Perez All-Weather Sky Model that will provide exact positioning of sunlight along with the skylight irradiance. The core of the algorithm will use the stochastic Next Event Estimation (NEE) Monte Carlo path tracer that will allow sampling separate components of illumination such as direct solar beams and diffuse scatterings independently. Overall, this approach provides curators with a non-destructive sandbox environment that enables instant dosimetric measurements without any drop in performance at approximately 60 fps rate.

1.4. Thesis Structure

Subsequently, the body of this paper is arranged in such a way that all the necessary information regarding ChronoLux will be delivered, from theory to experimental confirmation.

In Chapter 2, the main goals of this study will be described, namely, both interactive goals and metrological ones. Chapter 3 represents an extended review of the bibliography and describes recent achievements in the areas of physical dosimetry, offline ray tracing, and hardware accelerated global illumination. Chapter 4 contains the analysis of the theory of light transport, including the explanation of Lambert's Cosine Law, and Monte Carlo integrals. Chapter 5 describes the process of creation of the Digital Twin and focuses on the Texture Space Ray Tracing algorithm, as well as Next Event Estimation kernel. In Chapter 6, the experimental validation and testing of the proposed

framework is described in detail, including different conservation case studies and proving the mathematical consistency of the Monte Carlo convergence. In addition, Chapter 7 presents an elaborate user guide to using the software, and in Chapter 8 - conclusions and future perspectives are stated.

Chapter 2. Project Objectives

The main focus of this thesis is to create a hardware-accelerated Digital Twin, called *ChronoLux*, which will provide an interactive and physically-based lighting simulation for the conservation of cultural heritage. This solution is designed to fill the important missing link between computationally-heavy offline rendering systems (such as Radiance), which can be proven mathematically, and real-time IoT sensing devices.

In order to achieve this high-level goal and provide the necessary combination of interactivity and scientific accuracy, the development process is carried out based on certain technical and metrological goals. The high-level architecture related to these goals is outlined in Figure 2.1.

2.1. Technical and Architectural Objectives

- 01. Texture Space Ray Tracing Integration:** To address the drawbacks associated with traditional screen space algorithms, such as frustum culling that results in critical out-of-screen lighting information being removed from calculations, a special compute shader for DirectX 12 Raytracing (DXR) would have to be used to launch rays directly from the unwrapped UV texture map of 3D models in order to permanently bake the cumulative radiometric energy into all of its geometric surface regardless of the location of the observer.
- 02. Physical Environmental Modeling:** In order to combine the Perez All-Weather Sky Model with a special Next Event Estimation (NEE) compute algorithm to create a scientifically rigorous method to decouple the direct beams of sunlight from the diffuse ambient lighting, thus generating physically correct global illumination without loss of real-time processing ability.
- 03. Asynchronous Compute Architecture:** For a proper separation of concerns between the user interface layer and computationally demanding GPU path tracing engine, the program should allow users to modify physical properties (transmission coefficients of window glazing or wall reflectance, for instance) via an interactive UXML/USS user interface without breaking the continuity of the 60 FPS viewport interaction.

2.2. Metrological and Validation Objectives

- 04. Headless Statistical Telemetry:** In this study, an automated approach has been developed for virtual point-probe Lux sensors which operate purely in the background. The sensors are created for the purpose of computing spatial variance (σ^2), maximum exposure dose, and standard Monte Carlo convergence error boundaries while not interfering with the primary rendering process.

O5. Empirical Scenario Validation: Explicit validation of the above approach is achieved by means of running simulations of an entire year of meteorology (i.e., 8,760 hours of continual processing) for various preventive conservation scenarios. It is essential for establishing the efficacy of dynamic environmental control actions like the application of UV-absorbent windows or optimizing geometric occlusion of susceptible artworks.



Figure 2.1: High-Level Concept Architecture of the ChronoLux framework, illustrating the pipeline from raw heritage asset ingestion to real-time dosimetric data export.

The subsequent chapters of this thesis will methodically detail the design, theoretical foundation, programmatic implementation, and scientific validation of these stated objectives.

Chapter 3. Bibliographic Research

The creation of a hardware-accelerated Digital Twin for online light dosimetry is the result of the intersection of several complex scientific fields. The following section will provide an in-depth analysis of the literature related to preventive conservation metrology, mathematics of global illumination, and modern graphics processing technology. By analyzing the various approaches used today, the theoretical basis of this thesis will be defined, together with the limitations of current methods to be overcome by *ChronoLux*.

3.1. Preventive Conservation and Empirical Dosimetry

The need for light dosimetry is necessitated by the photochemically reactive nature of heritage objects. According to Michalski [1], both visible light and ultraviolet radiation are responsible for causing permanent molecular decay in natural and synthetic materials. In terms of photon energy, there is an inverse relationship between wavelength and energy level. The quantum energy of ultraviolet light, which occurs at wavelengths between 10 and 400 nm, is enough to break covalent bonds between pigment molecules, binding agents, and structural fibers.

Since light-induced deterioration is entirely accumulative, healing or reversing any damage is impossible simply by putting the artwork in darkness because the bonds have already been broken at a molecular level. To mitigate such problems effectively, the CIE set rigorous conservation limits and classified materials based on their varying degrees of vulnerability to light [2].

- **Insensitive Materials:** Artifacts are sensitive to exposure to light to differing extents. Materials like metals, stone, and ceramics are not vulnerable to photodegradation and, hence, have no special lighting requirements.
- **Low Sensitivity Materials:** Oil and tempera paintings, frescoes, and wood used in structures are capable of tolerating moderate exposure, up to 600,000 Lux-Hours per year.
- **High Sensitivity Materials:** Water colors, pastels, tapestries, and old papers, on the other hand, are extremely delicate and can only tolerate 50,000 Lux-Hours exposure in a year.
- **Extreme Sensitivity Materials:** Silk, certain unstable dyes, and early photos must be protected against all dangers posed by excess light, which is usually restricted to 15,000 Lux-Hours annually.

Going beyond these levels will hasten the process of entropy of artifacts, making it mandatory to establish accurate forecasting processes.

Traditionally, the approach that has been used extensively for evaluating such exposure is the measurement of incident light intensity. According to Silva et al. [3], the use of affordable data loggers to measure the microclimate in these locations is emphasized. In addition to this, Chianese et al. [4] describe the concept of “Internet of Cultural

Things” by suggesting the presence of dense IoT-based sensors within museums.

Discussion: While empirical sensor networks provide irrefutable real-world data, there exist several operational constraints associated with such systems that significantly limit their potential for use in modern conservation practices. First, it should be noted that these devices are purely reactive, with a physical sensor sending a signal when light interaction occurs, thus not enabling any preventive action before an exhibit is opened to light exposure. Second, the accurate simulation of the spatial distribution of irradiance at a huge three-dimensional surface would require installing an unrealistic amount of physical sensors on the subject. For instance, using thousands of sensors to cover a Renaissance sculpture would be both costly and visually unappealing. It could also be dangerous due to the fragility of such objects. As a result, to shift from reaction to prevention and simulation, mathematical models are required.

3.2. The Metrology of Global Illumination

3.2.1. Rasterization vs. Ray Tracing in Scientific Simulation

The visualization of virtual worlds in computer graphics has traditionally been implemented using rasterization algorithms. Rasterization maps 3D polygons to the 2D surface of the screen, computing their visibility through depth buffering. While extremely fast and highly suited for interactive purposes like games, rasterization is essentially incapable of scientific dosimetry. As the algorithm only takes into account direct visibility of objects (frustum culling), it heavily relies on ad hoc methods to create indirect lighting, including Screen Space Ambient Occlusion (SSAO), shadow mapping, and baked light probes, introducing significant bias.

Ray tracing, on the other hand, explicitly traces the actual path taken by the photons as they emanate from the light source, reflect from the geometry based on its material properties, and refract if necessary. Traditionally used for producing photorealistic images, ray tracing is the de facto standard for creating visual effects (VFX) for films and, most recently, top-tier video games due to the ability to trace perfectly accurate mirror reflections, soft shadows, and transparent materials’ refractions. In both of the cases, however, the goal remains that of deceiving human eyes and presenting realism without paying too much attention to the physical aspects in order to reduce rendering time.

Conversely, in order to achieve scientifically accurate dosimetry, one would need to simulate purely physical behavior of light with no regard to any virtual cameras or human perception involved. For that purpose, simulations need to be performed using unbiased and purely mathematical ray tracing which will allow the traced photons to serve as purely radiometric data. This technique is mathematically provable and thus represents the only possible algorithm for the measurement of the actual irradiation and Lux-Hour accumulation in a virtual world.

3.2.2. The Rendering Equation and Monte Carlo Integration

The key model that forms the backbone of physically based ray tracing is known as the Rendering Equation and was first introduced by Kajiya [5]. It represents a balance between the amount of radiance leaving a surface point through all directions, and the

summation of two quantities – radiance coming from emitting objects and scattered radiance from the hemisphere around the surface point.

Due to the fact that solving this infinite-dimensional equation is impossible for any graphics engine, an approximate method called Monte Carlo integration must be used. It entails sampling a random subset of rays in order to estimate the outcome of the integral. In other words, the more samples used, the higher the precision of the estimate – a principle defined by the Law of Large Numbers.

However, uniformly sampling random rays proves to be rather wasteful because many of them will not interact with any source of illumination, hence producing a noisy picture with high variance. In order to overcome this issue, the importance sampling technique can be applied. This approach uses a probability density function (PDF) that defines where rays should be cast based on the expected light distribution (e.g., cosine-weighted importance sampling for Lambertian surfaces).

In case the environment is enclosed and bounded, stochastic termination methods, like Russian Roulette, need to be applied in order to avoid infinite bouncing of rays. Instead of limiting their path length with an artificially chosen number of bounces, Russian Roulette ends the path when encountering dark surfaces which absorb more light.

3.2.3. Physical Material Simulation and Energy Conservation

Fidelity of the illumination is based on the rigour of material definitions. When a photon hits a geometrical surface, the physics behind what happens next can be described by the surface's Bidirectional Reflectance Distribution Function (BRDF) or Bidirectional Transmittance Distribution Function (BTDF). The combined effect describes the distribution function for the way the irradiance is reflected, absorbed, and transmitted through the hemisphere defined by the surface normal.

Modern computer graphics often use elaborate PBR materials that can be described by the state-of-the-art microfacet models such as the GGX distribution. According to Pharr et al. [6], these microfacet equations are employed to simulate the fine details of surface roughness and metals to achieve photorealistic visuals. However, their complex nature comes at an enormous computational cost while making a relatively minor contribution to the estimation of radiometric energy density on flat architectural surfaces.

ChronoLux does not use complex visuals of modern PBR materials but focuses on thermodynamically rigorous calculations. In metrological terms, according to the First Law of Thermodynamics (Energy Conservation), the sum of Reflectance (R), Transmittance (T), and Absorptance (A) of any material must be equal to one ($R + T + A = 1$). Thus, the reflectance of any hypothetical material reflecting 60% and transmitting 30% of the incident irradiation, like a tinted window glass, must have 10% absorptance.

Correctly accounting for these factors is critical for preventive conservation purposes. Walls are normally Lambertian, which means the photon is equally distributed in all directions after interaction with the wall surface, affecting the level of indirect illumination in the room. By contrast, windows imply complex interactions of photons and materials including transmission and refraction. A correct estimation of the transmittance graph of the UV-blocking glass or reflectance factor of polished marble flooring is crucial as it determines how the Lux-Hour dosage will accumulate on the artifacts.

3.2.4. Limitations of Offline Rendering

The foremost example of these approaches in architecture and conservation is the Radiance program [7]. The Radiance approach combines backwards ray tracing together with irradiance caching to solve the Rendering Equation with extremely high metrological precision.

Discussion: In contrast to other approaches to rendering, Radiance allows for the creation of lighting distributions that are provably correct from a mathematical standpoint. For this reason, it has become the gold standard for offline dosimetry calculations. The reliance on CPU bound processing however leads to significant inefficiencies. It takes several hours for the calculation of a full year of exposures, thus ruling out any interactive examination of the potential conservation strategies.

One of the alternatives to this method, Precomputed Radiance Transfer (PRT) [8], aims at solving the problem of real-time global illumination. PRT is achieved through precomputation of light transport into spherical harmonics. However, its applicability suffers greatly when the environmental effects are not static. An adjustment of the object orientation or transmittance properties of window glazing requires recalculating the entire PRT lightmap.

3.3. Hardware Acceleration and Heritage Digital Twins

The creation of highly accurate dosimetry with the capability of real-time interaction requires architectural thinking to be updated. As stated by Jouan & Hallot [9] and Niccolucci et al. [10], the concept of "Heritage Digital Twin" goes beyond that of traditional 3D representation or, otherwise known as HBIM, and becomes a dynamic replication continuously fed with data from reality to make decisions about preventive conservation.

3.3.1. Bounding Volume Hierarchies and Ray Traversal

The introduction of the Rendering Equation in this context at 60 Hz renders a need for complex spatial acceleration structures. Finding an intersection of a given ray with several million raw geometric triangles mathematically becomes impossible to compute within seconds. For that reason, scenes have to be split into Bounding Volume Hierarchy (BVH).

BVH is a kind of tree-like structuring where all geometries are placed into Axis-Aligned Bounding Boxes (AABB). The root of such structure contains the whole scene and, consequently, all geometries are successively divided by nodes until leafs, containing actual triangles, are reached. When tracing rays, we are checking intersections with bounding boxes. In case a ray misses a parent AABB, the algorithm ignores all of those millions of triangles turning an $\mathcal{O}(n)$ problem into a logarithmic $\mathcal{O}(\log n)$ search.

Surface Area Heuristi (SAH) is utilized to construct these trees optimally such that it reduces the probability of any intersection between rays and bounding boxes by assessing the cost of node traversal and triangle intersection.

3.3.2. DirectX Raytracing (DXR)

Traditionally, BVH traversal and construction was carried out using software on CPUs. In the words of Burgess [11], this approach was turned upside down in terms

of the NVIDIA Turing GPU architecture that came equipped with hardware silicon RT Cores for accelerating BVH traversal and triangle intersection.

The DXR API provides this functionality to developers by offering two acceleration structures for managing static and dynamic objects respectively. The first of these is the Bottom Level Acceleration Structure (BLAS). The bottom level structure contains all the raw triangle information of each mesh. Although extremely optimized, its creation requires high computational effort due to the complexity involved in the process. In contrast, the second acceleration structure—the Top Level Acceleration Structure (TLAS)—contains light instance references to BLAS that allow transformations to be applied instantly without rebuilding.

DXR provides an updated shader pipeline that makes use of a wide array of frameworks like ChronoLux:

- **RayGeneration Shader:** The first shader program used in the pipeline, responsible for generating primary rays. Unlike the traditional method of launching rays via camera, this process launches rays in Texture Space.
- **ClosestHit Shader:** Triggered upon successful intersection between a ray and solid geometry. At this point, it evaluates properties like BRDF, physics albedo, and diffuse interreflection.
- **AnyHit Shader:** Executed before Closest Hit to evaluate opacity or perform alpha test. This shader is required in scientific dosimetry to accurately calculate physics transmissivity across refractive window glazing without stopping the shadow ray.
- **Miss Shader:** Executed once the ray travels past local geometry. When working on environmental simulations, Miss shaders evaluate global illumination techniques such as Perez All-Weather Sky Model [12] to accurately compute solar irradiance.

3.3.3. Variance Reduction and Texture Space Rendering

For optimal performance of Monte Carlo Integration, Veach and Guibas [13] provided the necessary mathematical foundations for incorporating sampling algorithms that enabled the utilization of Next Event Estimation (NEE). NEE greatly helps to reduce the problem of variance through explicit casting of a shadow ray to the main light source in each bounce. The technique ensures fast convergence, even when there is indirect lighting, since it disassociates sunlight rays from ambient ones.

Moreover, for isolating fast ray traversal from the boundaries of the viewport of a camera, Munkberg et al. [14] presented modern texture-space cache methods. Traditional screen space renderers throw away useful lighting information that falls out of the current frustum view. The new technique renders images in texture space by shading objects in their corresponding UV mappings regardless of camera occlusions.

3.3.4. Solar Telemetry and Ephemeris Models

A precise simulation, especially one involving an entire meteorological year consisting of 8,760 hours, demands the exact position of the Sun in the sky. The solar position determines the angle of incidence of radiation, thus the shadows cast on any objects, and the amount of transmitted light through windows.

This can be obtained using various mathematical models of solar position from the scientific computing libraries, known as Solar Ephemeris. For instance, there exists a library called Solar Position Algorithm (SPA) formulated by Reda and Andreas [15]. This library calculates the solar zenith and azimuth positions based on latitude and longitude

position of the heritage sites, together with the Julian day and local time at any given point.

Apart from precision, calculation of the solar position is vital for computational optimization. Path tracing is quite computational intensive. Calculation of solar zenith position prior to the execution of the path tracing process allows the simulation to know if the Sun is below the horizon line. In case when it is night time (time between sunset and sunrise), then the process of path-tracing for the specific hour is not necessary, therefore reducing computation time significantly.

3.3.5. The Perez All-Weather Sky Model

The correctness of the virtual environment dose estimation model depends on the accuracy of the external light source driving it. In traditional three-dimensional rendering, the simplest form of light direction is usually used to simulate sunlight and set the sky color to uniform color; however, the physical process of atmospheric light scattering has some significant complexity. The complexity is solved by the use of scientific metrology through precise atmospheric math calculations, namely, the Perez all-weather sky model [12].

The Perez model is the empirical mathematical model that calculates the relative luminance of any sky point with regard to real meteorology conditions. Opposed to other simple models describing only ideal conditions of absolutely sunny or cloudy weather, Perez uses two atmospheric parameters:

- **Clearness (ϵ):** An index used to describe the ratio between direct solar radiation and diffuse sky radiation, giving a measure of how much clouds there are and their thickness.
- **Brightness (Δ):** An index that measures atmospheric luminous opacity, dependent on aerosol concentration and humidity.

With all these parameters taken into account along with precise solar angle and azimuth values, the model creates a very precise continuous luminance distribution. Such an approach incorporates modeling of two important scattering effects within the atmosphere, namely, circumsolar brightness, which results from Mie scattering near the sun and horizon brightness which is a result of the luminance gradient formed due to Rayleigh scattering in denser atmospheric layers around the horizon. In terms of the ChronoLux system, once the path traced ray leaves the building using DXR Miss shader, its vector direction is used to obtain a physical spectral radiance value for this particular micro angle within the sky dome.

Conclusion: The approaches examined within the current chapter demonstrate that even though offline rendering tools such as Radiance give precise results, and IoT sensors provide actual information, these methods cannot predict real-time changes and thus require interaction. Through a combination of the precise calculations of Monte Carlo path tracing algorithm and NEE technique suggested by Veach [13] along with the Perez sky model [12] and powerful hardware acceleration obtained through use of DXR [11] and BVH traversal along with texture space ray tracing [14], the proposed approach will fill this gap in preventive conservation.

Chapter 4. Analysis and Theoretical Foundation

In this chapter, a clear definition of theoretical algorithms, mathematical abstractions, and overall logic behind the *ChronoLux* design is attempted. Following the literature review, it becomes apparent that current methods of simulation suffer from significant shortcomings that render them unsuitable for application within interactive heritage metrology. Subsequently, in the following section, the theoretical measures suggested to tackle the aforementioned problems are outlined, focusing purely on abstract data manipulation and physics of light transfer.

4.1. High-Level Theoretical Pipeline

The *ChronoLux* pipeline is an endless sequence of mathematical computations that turns abstract three-dimensional data into radiometric energy build-up. From an architectural perspective, the *ChronoLux* system can be broken down into three separate domains. These include the Initialization Domain, Radiometric Integration Domain, and Temporal Accumulation Domain. As described earlier in Figure 2.1, *ChronoLux* operates on a closed-loop basis, using time-discrete iterations for its operations.

4.1.1. The Initialization Domain

Before the execution of radiometry calculations, a physical definition of the virtual environment is required. This initialization starts from the import of geometrical mesh information. From the theoretical standpoint, each artifact consists of a discrete collection of interconnected triangles forming a manifold geometry. Each of these triangles is connected to physically realistic material behavior, described by the strictly conservative thermodynamic equation ($R + T + A = 1.0$) outlined in Chapter 3.

The geographical and temporal status of the simulation is established at the same time. According to the theoretical description, the coordinates of the geographic latitude (Φ), geographic longitude (λ), geographic timezone offset, and the initial Julian Day of the physical heritage location are necessary inputs into the astronomy ephemeris calculation algorithm. After the completion of the first stage, a Bounding Volume Hierarchy (BVH) with a statistical optimization applied to the spatial information, as well as the solar zenith angle (θ_z), and solar azimuth angle (γ_s) for the current temporal step (t) are derived. If the calculated solar zenith angle is less than zero ($\theta_z < 0$), then, due to the system's automatic culling mechanism, the radiometric calculation process does not have to be executed at all, moving to the next temporal iteration point.

4.1.2. The Radiometric Integration Domain

After both sets of parameters have been initialized, the system proceeds to the next stage, which is referred to as the Radiometric Integration Domain. This is where the primary mathematical algorithm for calculating irradiance takes place. Unlike in conventional rendering techniques, which use the projection of the view frustum onto the two-dimensional space, all computations take place in Texture Space, an idea that will be further explained in Section 4.2.

In this domain, a computational solution is sought for the Rendering Equation at all acceptable points of the artifact's surface. Rays are propagated stochastically through the hemisphere to calculate both direct sunlight and diffuse ambient light using the technique of Next Event Estimation (NEE). Directions of escaping rays are compared to the Perez All-Weather Sky Model, allowing the system to receive exact spectral values of radiance. At the end of computations, a two-dimensional grid of data that corresponds to instantaneous physical irradiance (in Lux) is obtained.

4.1.3. The Temporal Accumulation Domain

The last theoretical step comprises the constant temporal summation of the irradiance values at each moment of time. Photochemical decomposition is always an accumulative process, so the calculated radiometric information will be reused rather than thrown out after one-time processing. In particular, the instantaneous illuminance (Lux) value is multiplied by the temporal duration of the current simulation step (Δt) to derive an accumulative dosage in Lux-Hours. The accumulative dosages will be added up using a persistent floating point array. After that, the temporal counter ($t = t + \Delta t$) is advanced, and the procedure is iterated until the end of the chosen period of time (e.g., meteorological year).

4.2. The Texture Space Ray Tracing Paradigm

The first major theory in which *ChronoLux* deviates from regular computer graphics is that of using the Texture Space Ray Tracing method of computation. To illustrate why such a computational method is needed in science metrology, its advantages, as well as the mathematical limitations of conventional methods, will be demonstrated below.

4.2.1. The Mathematical Flaws of Screen Space Metrology

First, in a three-dimensional engine, there is always a virtual viewer defined by a set of parameters including a field of view, distance from the screen, and near/far planes. During rendering, a two-dimensional pixel grid is used. For each pixel, there is a ray traced from the camera's optical point through the pixel into three-dimensional space.

Such an approach is mathematically exact, but it is flawed when we need to accumulate some radiation on all geometric surfaces. First, the problem is related to occlusion and perspective projection. When tracing rays from the virtual viewer, it becomes clear that the system will trace those rays that actually intersect the polygons. Any surface turned away from the viewer will not be computed because it simply will not be covered by any rays emanating from the viewer. It is obvious that if a statue is turned around the wrong way, no radiation can be accumulated on its posterior surface; the same holds true for other surfaces.

4.2.2. Texture Space Mapping

To accomplish a full separation between the simulation and the artificial restrictions of camera placement, *ChronoLux* uses Texture Space Mapping. In 3D geometry, each vertex of the polygon corresponds to a point on the 2D UV plane through (u, v) coordinates in three-dimensional geometry, which is called the UV map.

The process of traversing the UV map in the *ChronoLux* theory does not involve traversing the pixels on the computer screen but involves traversing the discrete texels on the UV map. This allows for the recreation of geometry in world space by calculating P (position) and N (surface normal).

Let the UV map be a 2D matrix M of size $W \times H$. For every element $M_{i,j}$:

1. Read the pre-computed world-space position vector $P = (P_x, P_y, P_z)$ corresponding to the (u, v) coordinate.
2. Read the pre-computed world-space normal vector $N = (N_x, N_y, N_z)$.
3. Use P as the origin point for the dosimetric path tracer.
4. Use N to orient the sampling hemisphere.

This method ensures that each infinitesimally small patch of the surface of the object is evaluated simultaneously, without regard to the virtual camera orientation. Global consistency of the physical light transport throughout the entire structure of the heritage object is guaranteed.

4.3. Dosimetric Path Tracing Algorithm

For the sake of simplicity, the origin (P) and the direction (N) of the rays as established using the Texture Space technique are known; hence, it is necessary to solve the global illumination equation in order to find the exact amount of energy hitting the particular point. This task can be accomplished using a highly efficient Monte Carlo algorithm with Next Event Estimation (NEE).

4.3.1. Stochastic Hemisphere Sampling and Energy Attenuation

Firstly, to find the total irradiance on a surface, the sampling directions have to be found, which involves generating directions distributed evenly over the upper hemisphere centered at the point where the sampling takes place. Within the framework of the Monte Carlo algorithm, such distribution of samples can be represented using a stochastic directional vector ω_i .

To improve convergence rates and minimize variance, *ChronoLux* uses Cosine-Weighted Importance Sampling. Instead of uniformly sampling random vectors, the probability density function used here tilts towards the vertical axis that passes through the surface normal. This method mirrors the physics behind lambertian diffuse surfaces when light coming from the vicinity of the surface normal carries more energy (proportional to $\cos(\theta)$) than that coming in from a large angle θ .

As these stochastic rays traverse the digital scene, they interact with other surfaces. The core principle behind this interaction is energy absorption, meaning the photonic path will dissipate its energy after each encounter with another surface. Thus, an object placed in a room with dark mahogany walls (with high Absorptance, A) will reflect indirect rays with significantly less energy than those reflected from white plaster (high Reflectance, R).

This effect can be easily modelled mathematically through the use of a throughput multiplier. Starting with an initial value of 1.0, which means perfect energy transmission, throughput gets multiplied by Reflectance R at each encounter. As $R < 1.0$ for any physical surface (following the law of conservation of energy $R + T + A = 1.0$), the ray loses energy at an exponential rate after each reflection.

In order to prevent simulation from becoming infinitely long, because of an endless number of bounces in enclosed space, a finite bounce depth limit should be set. For example, typical indirect path lengths do not exceed 3 or 4 bounces since additional reflections carry too little residual energy to matter further.

4.3.2. The Physics of Refractive Transmission

A basic element of heritage metrology involves evaluating the impact of external light that passes through physical obstacles, especially the effects of window glazing. This requires realistic modeling of refraction for light.

When a path traced ray meets a boundary between two media (say, a pane of glass), the ray does not necessarily stop or bounce off. Rather, it changes direction based on the IORs of the two media (air-glass) according to Snell's Law:

$$\frac{\sin \theta_1}{\sin \theta_2} = \frac{n_2}{n_1} \quad (4.1)$$

In the above-mentioned framework, n_1 and n_2 stand for the refractive indices of the media from which the ray starts and where it reaches, respectively; while θ_1 and θ_2 stand for the angles of incidence and refraction, respectively, regarding the normal to the interface.

As far as the *ChronoLux* abstract model is concerned, the transmission vector is calculated via Snell's Law whenever the beam penetrates a transparent body, allowing the ray to proceed with its travel through the solid glass. At the same time, the throughput of the ray is adjusted with respect to the particular material's Transmittance (T) factor. If a conservator applies UV-blocking film with a Transmittance of 0.45, the ray throughput is immediately decreased by 55%, thus providing an adequate simulation of the real-world effect produced by the intervention.

4.3.3. Next Event Estimation (NEE)

In a pure Monte Carlo path tracer, all sampling is random, and the tracing algorithm can use only random reflection from surfaces. The chance of finding the sun's small angular diameter through random reflections inside a dark museum is practically nil; therefore, there is significant variance in the simulation, and one would need to calculate millions of samples to converge toward the correct answer, rendering the calculation unfeasible at least with respect to the 8,760-hour simulation.

To alleviate this problem, *ChronoLux* employs the technique known as Next Event Estimation (NEE).

At every single interaction point P across the entire bounce sequence, the algorithm performs the following logic:

1. **Direct Shadow Ray:** A deterministic vector ω_{sun} is generated with the direction towards the exact current location of the sun. "Shadow ray" tracing occurs along ω_{sun} . If this ray leaves the bounded region without intersection with any opaque objects, the direct irradiance is calculated using the Perez model and added to the total dose. In case there is an intersection with a transparent object (such as

window glass), the light intensity will be reduced by a factor of T which is obtained using Snell’s law.

2. **Indirect Bounce Ray:** At the same time, a stochastic vector ω_i is generated using Cosine-Weighted Importance Sampling. Tracing is performed along this vector to calculate the indirect light caused by reflection from walls and floors.

Through its clear disentanglement of the deterministic solar irradiance and the random scattering of the ambient sky dome, NEE guarantees that the massive energy content of the sun is properly accounted for at every stage of the simulation, thus reducing the number of samples needed by many orders of magnitude and ensuring metrological stability. The total irradiance of the starting surface point is thus computed as the sum of all successful shadow rays cast from the initial surface point through multi-bounces.

4.3.4. Miss Evaluation and Perez Integration

The last piece of the integration domain puzzle is that of Miss evaluation. When the stochastic ray (be it the direct shadow ray or the indirect bounced ray) finally escapes the digital twin architecture, it will need to retrieve the radiance value from the physical world surrounding the building.

Contrary to returning a hardcoded color, the escaping ray direction, ω_{escape} , is mathematically inserted into the Perez All-Weather Sky Model function. This function calculates the elevation (θ_z) and azimuth (γ_s) of ω_{escape} against the current solar location (θ_z, γ_s) and uses the meteorological data of Clearness (ϵ) and Brightness (Δ) to calculate the physical luminance value of ω_{escape} . If ω_{escape} happens to be near the solar disk, the result will undergo an exponential increase to simulate the Mie circumsolar scattering effect. Otherwise, if ω_{escape} is close to the horizon, then the Perez output transitions smoothly to simulate Rayleigh scattering.

4.4. The Parallel Determinism Problem

One of the problems that arises when moving from the abstract realm of mathematics into practical implementations in highly parallel architectures, such as today’s GPUs, is non-determinism in floating-point additions.

A typical parallel tracing algorithm traces hundreds of millions of rays independently using thousands of threads simultaneously. Once the rays finish traversing the scene, all threads attempt to write their irradiance into a shared matrix of accumulations M .

As opposed to integer numbers, floating-point operations are non-associative, meaning that $(A+B)+C \neq A+(B+C)$ because of tiny differences that occur during precision rounds in IEEE 754. As a result, the final value of luminance will be different each time the algorithm runs if the order of writing to the accumulation matrix M by the threads is random.

In video games, this minute difference is not detectable. When applied to science, however, non-determinism becomes unacceptable since an analysis tool must return the same results every time a test simulation is conducted.

The solution to the problem for the *ChronoLux* framework involves making sure that ray tracing is done and accumulations happen in a predetermined manner, through mathematical structuring of geometry or seed points. This solution may add minimal computational expense, but it guarantees mathematical consistency and reproducibility.

4.5. Mathematical Dose Accumulation

The key theoretical objective of the *ChronoLux* tool is not only creating a photographically accurate representation of light but rather finding the overall photochemical risk posed over a long period of time. It implies moving from considering an instantaneous value of illumination to a cumulative dosimetry one.

4.5.1. Temporal Integration Logic

Let $E_t(P)$ be the instantaneous physical irradiance (measured in Lux) estimated with the help of dosimetric path tracer on a particular surface point P in a particular moment of time t .

Light illumination in a physical museum is a continuous process, while a digital simulation has to discretize time into particular segments (Δt). The choice of Δt significantly affects not only the accuracy of the calculation but also its performance. A traditional meteorological simulation usually takes 1 hour for each Δt . It allows tracking a motion of the Sun and cloud cover change in 8,760 steps within a typical year.

Cumulative dose $D(P)$ delivered to a particular surface point P for the whole period of time T is theoretically equal to the following sum:

$$D(P) = \sum_{t=0}^T E_t(P) \cdot \Delta t \quad (4.2)$$

Where:

- $D(P)$ is the total accumulated exposure in Lux-Hours.
- $E_t(P)$ is the calculated irradiance in Lux at time step t .
- Δt is the duration of the time step in hours.

4.5.2. Floating-Point Precision and Heatmap Generation

In order to retain the accumulating radiation dosage $D(P)$, the program uses the previously discussed Texture Space M matrix. Standard bitmap file formats (such as the 8-bit PNG or JPEG) cannot be used to store dosimetric data, as they can only store values from the set of integers ranging from 0 to 255. An extremely sensitive piece of art may amass as much as 50,000 Lux-Hours during its lifetime, which will quickly surpass the capacity of the 8-bit variable.

Therefore, according to the theory, the system makes use of a 32-bit floating point data structure (RFloat). A 32-bit float can theoretically encode 3.4×10^{38} , providing practically infinite precision for unbounded accumulation.

As the simulation advances within the temporal dimension, the matrix elements continuously increase in value. The abstract mathematical information is then mapped to a heatmap for visual analytical representation of areas where photosensitivity risks are critical.

4.6. Conclusion

Thus, in conclusion, the theories described in the current chapter make *ChronoLux* a purely metrological simulation software. Abandoning the bias-laden approach of the Screen Space rendering in favor of the Texture Space ray tracing method guarantees

complete geometrical coverage. Furthermore, Next Event Estimation, together with the Perez Sky Model algorithm, ensures exact computation of the light transport. Finally, accumulating irradiation in an arbitrarily precise floating point matrix allows for accurate simulation of photochemical degradation for meteorological years.

Chapter 5. Detailed Design and Implementation

In this chapter, a thorough explanation is provided regarding the software architecture, the accelerated algorithms running on hardware, and the rigorous mathematical systems used in order to create a ChronoLux digital twin. The practical realization of radiometric and conservation physics theories on high-performance desktop software requires a highly sophisticated coordination between CPU scheduling, memory transfers, and extensive GPU parallelism.

The historic inability to provide real-time interaction with the predictive workflow using the CIE and Michalski principles for preventive conservation can be attributed to the lack of interactive offline software, such as Radiance, that requires a lengthy amount of CPU computations in order to simulate only one temporal frame. ChronoLux solves this problem by creating a hybrid architectural solution, using the Unity 3D engine and the DirectX 12 graphics API in order to allow the usage of hardware-accelerated DirectX Raytracing (DXR). Each component of the computational pipeline from mathematical astrometry to GPU memory management and asynchronous telemetry extraction will be analyzed.

5.1. Software Architecture Overview

In order to run sufficient amount of ray queries for a yearly dosimetric simulation without causing problems with system operation, a strictly decoupled architecture was chosen. The software stack is divided into two major domains: the first domain contains the C# CPU side of the program, including the UI, file I/O, deterministic structural sorting and mathematical astrometry and the second one contains the HLSL compute shaders running the highly parallelized Monte Carlo integrations on the GPU.

Data flow is tightly controlled. The project parameters are serialized to JSON with the help of `ChronoProject` object and parsed with the `ProjectManager` that feeds the `LightDoseSimulator`. The `LightDoseSimulator` acts as an orchestrator controlling the time and dispatching ray queries to the DXR kernels.

5.2. Core Simulation Engine (C# CPU Orchestration)

The central logic should guarantee mathematical determinism, avoid hardware starvation, and respond to user interactions.

As one can see on the diagram above, the `LightDoseSimulator` depends on multiple important subsystems: `SunCalculator` for mathematical astrometry, `IrradianceBaker` for hardware dispatch and `RuntimeModelLoader` for asynchronous file I/O.




Figure 5.1: High-level simulation pipeline demonstrating the decoupling of CPU orchestration from GPU execution.




Figure 5.2: Overall Architectural System design of *ChronoLux*, detailing data paths between the Engine, the GUI, and the OS filesystem.




Figure 5.3: C# CPU Class Architecture and Data Flow orchestration within *ChronoLux*.

5.2.1. The Chronological Engine and Time Stepping

The primary job of the `LightDoseSimulator` is to increase the time by Δt and to calculate the precise solar altitude and azimuth for the Perez sky model [12]. The simulation should calculate one meteorological year without causing the Timeout Detection and Recovery (TDR) event, that might forcefully stop graphics drivers operation after the certain timeout. Such event occurs after a roughly two second timeout in the certain task and is handled with the help of the C# `IEnumerator` Coroutine called `RunSimulationInternal`. As detailed in Appendix 9.2 (Listing 9.3), the time stepping loop is fully dynamic and goes over the range from `startDay` till `endDay`.

For each day, the coroutine uses the `SunCalculator` to calculate the exact local sunrise and sunset boundaries based on user's defined latitude, longitude and UTC offset and thus, avoiding the unnecessary ray casting during night hours. In case of equal sunrise and sunset time (a rare case in polar nights), the current day is simply skipped.

Aggressive computational culling is used in the hourly step loop. The engine ignores all steps where both calculations of direct-beam lux and diffuse-sky lux are negligible (lower than 0.1 Lux). Such optimization allows us to save millions of useless Monte Carlo ray casts per simulated year.

During the active simulation period, the coroutine explicitly uses the `yield return null;` instruction in order to put itself to sleep until the next engine frame. Such approach allows to give a thread back to the Unity main loop thus maintaining responsive UI, proper animation of progress bar and satisfying TDR watchdog.

5.2.2. Mathematical Astrometry and Solar Ephemeris

In order to get metrologically correct results, the solar positioning is calculated at each hour. The `SunCalculator.cs` class (see Listing 9.5) performs rigorous mathematical models [16] in order to get solar disc vector position and atmospheric attenuation and thus, getting physically accurate `BeamLux` and `DiffuseLux` values.

First of all, the Julian Day (JD) epoch is calculated for the local simulation time

[17] and Julian centuries T since J2000.0 are derived:

$$L_0 = (280.46646 + T(36000.76983 + T \times 0.0003032)) \pmod{360^\circ} \quad (5.1)$$

$$M = 357.52911 + T(35999.05029 - 0.0001537 \times T) \quad (5.2)$$

To take into account the Earth orbit eccentricity, the equation of center C is calculated via Fourier series expansion and thus, giving us the apparent longitude λ_{app} :

$$C = \sin(M)(1.914602 - 0.004817T) + \sin(2M)(0.019993) + \sin(3M)(0.000289) \quad (5.3)$$

The Mean Obliquity of the Ecliptic ϵ_0 is corrected with the nutation correction and gives the Solar Declination δ , that is an angular displacement of the sun relatively to the celestial equator:

$$\delta = \arcsin(\sin(\epsilon) \sin(\lambda_{app})) \quad (5.4)$$

Knowing δ , the Equation of Time (E_{qT}) translates the solar time into the civil time. Together with terrestrial longitude, this information gives us True Solar Time (TST) and Hour Angle H :

$$H = \frac{TST}{4.0} - 180^\circ \quad (5.5)$$

The Solar Zenith Angle θ_z and Azimuth ϕ can be defined relatively to latitude Φ :

$$\cos(\theta_z) = \sin(\Phi) \sin(\delta) + \cos(\Phi) \cos(\delta) \cos(H) \quad (5.6)$$

Having the astrometric vector position, we calculate clear-sky direct-beam irradiance E_{beam} with the help of Kasten–Young formula for optical air mass AM [18]:

$$AM = \frac{1}{\sin(\alpha) + 0.50572(\alpha + 6.07995)^{-1.6364}} \quad (5.7)$$

$$E_{beam} = 127500 \times 0.7^{AM^{0.678}} \quad [\text{Lux}] \quad (5.8)$$

Where α denotes solar altitude angle ($90^\circ - \theta_z$) and 0.7 represents baseline atmospheric transmittance factor for clear-sky conditions. For the ambient lighting, we use a simple Haurwitz clear-sky diffuse model [19] that estimates the horizontal diffuse illuminance $E_{diffuse}$ relatively to the total solar constant and atmospheric penetration. Precise **BeamLux** and **DiffuseLux** values are provided to the HLSL compute kernel on each time step.

5.3. Scene Management and Hardware Interfaces

5.3.1. Asynchronous Runtime Mesh Parsing

ChronoLux is delivered as a compiled desktop metrology application and therefore cannot depend on pre-compiled Unity Scene assets. It should dynamically load arbitrary user provided `.obj` files at runtime from the local filesystem. Heritage artifacts, such as laser scanned church interior or high resolution photogrammetry models, could have millions of vertices. String-based parser working synchronously on the main Unity thread would make the application unresponsive and could cause the Operating System "Application Not Responding" failure. Therefore, runtime mesh parsing is implemented asynchronously.

As seen in Appendix 9.2 (Listing 9.4), the `RuntimeModelLoader.cs` resolves this memory bottleneck using a heavily parallelized chunk parser. File reading and string parsing of vertices, normals and UV coordinates are executed entirely by a separate `Task.Run()` background thread running on various logical cores of the CPU. As the parsing takes place in the background, the main thread constantly checks the condition `while (!parseTask.IsCompleted)`, calling `yield return null;` and updating the `ProgressBar` UI via `UpdateProgress()`. Once the enormous memory arrays are completely filled, control is passed back to the Unity main thread, which performs the final `Mesh.SetVertices()` VRAM upload and then creates the `GameObjects`. Such architecture ensures full UI responsiveness during the most data-intensive parts of the process.

5.3.2. Dynamic Material Serialization and Hierarchy Keying

Once the artifact is loaded with the runtime loader, all the original materials are stripped away completely. The simulation applies physically validated presets (such as Brick, Hardwood, Antique Glass) to the individual submeshes, and this assignment is stored securely in the `ChronoProject` JSON file.

In order to reconstruct these materials safely across various application sessions, the `RuntimeModelLoader.cs` implements a unique string-keying approach. Relying solely on the `GameObject` name is impractical because there are multiple components called, for example, `Wall` or `Window_01` in the `.obj` files. The `GetObjectKey` function creates a unique hierarchy key:

```

1 private string GetObjectKey(GameObject go) {
2     if (go.transform.parent != null && go.transform.parent != modelRoot
3     )
4         return $"{go.transform.parent.name}/{go.name}";
5     return go.name;
6 }
```

With this `RootName/MeshName` key combination, it is guaranteed that once the `ApplySavedMaterial` logic loops over the loaded mesh components, the geometry will be matched securely to the serialized `project.objectNames` array and the exact GUID of `project.materialPresetNames` will be found, thus preventing any material crossover errors.

5.3.3. Spatial Normalization

There is a huge variability in imported geometries - some architectural models are measured in millimeters, other ones in meters, and some geographical surveys in kilometers. In order to prevent floating-point precision losses in the DXR kernel - where the rays should intersect coordinates in wide ranges - *ChronoLux* employs a rigorous spatial normalization procedure through the `CenterAndScaleModel` function.

First, the global bounding box enclosing all the submeshes attached to the artifact root is determined and the maximum dimensional size is calculated:

$$MaxDim = \max(Width, \max(Height, Depth)) \quad (5.9)$$

Then, the uniform scaling factor is determined:

$$Factor = \frac{2.0}{MaxDim} \quad (\text{assuming } MaxDim > 0) \quad (5.10)$$

Every child vertex transform is translated in such a way that the mathematical center of the object is placed at the (0,0,0) and then uniformly scaled with the factor. Thus, the whole heritage site regardless of its original CAD scale is forced to occupy the strict 2.0 meters bounding box in World Space. Since the DXR raytracer works with 32-bit floating-point calculations, confining intersection geometry to the $[-1.0, 1.0]$ range around the origin maximizes IEEE 754 mantissa precision and prevents such problems as shadow acne or geometric clipping. Then, the user’s semantic transformations (position and scale offsets read from the **ChronoProject** JSON) are applied only to the normalized root object.

5.3.4. Deterministic BVH Building

As described in the theoretical radiometry literature, the Monte-Carlo integration is prone to floating-point associativity errors. An even greater problem in the hybrid CPU-GPU engine is the structural non-determinism. The Unity’s **FindObjectsByType()** API doesn’t provide any deterministic order guarantees of its return value for different OSs and CPU architectures. If the assignment of the ray’s geometric **InstanceID** and corresponding physical **SimulationMaterial** is varied from run to run, then the dosimetric simulation becomes scientifically non-reproducible.

In order to address this issue, the **IrradianceBaker.cs** initialization makes sure about the determinism through collection of all active **Renderer** components and their explicit sorting with highly granulated $O(n \log n)$ string-based comparison. Every submesh receives the unique **SortKey**:

$$\text{SortKey} = \text{ScenePath} | \text{TransformHierarchy} | \text{Type} | \text{MaterialSlot} \quad (5.11)$$

The **TransformHierarchy** is determined with the recursive traversal of the parent chain where the exact **GetSiblingIndex()** padded up to six digits (like 000002:Wall/000000:Window). This granulated sorting makes sure that the structural Index and Vertex buffers uploaded to the GPU along with the **_SimulationMaterials StructuredBuffer** have mathematically identical and well-aligned content, which ensures the reproducibility in the lux-hours accumulation.

5.4. Hardware-Accelerated Path Tracing (HLSL / DXR)

The primary computational bottleneck of *ChronoLux* resides in the ray tracing pipeline. For each pixel on the artifact’s UV map, stochastic rays are generated to calculate the rendering equation over long simulated periods of time [5].

5.4.1. Memory Management and DXR Acceleration Structures

DXR acceleration structure hierarchy consists of two tiers: the Bottom-Level Acceleration Structure (BLAS) storing the raw triangles data, and the Top-Level Acceleration Structure (TLAS) which defines instances of the BLAS in space. In *ChronoLux*, TLAS is built and managed with the help of Unity’s **RayTracingAccelerationStructure**.

Memory is carefully aligned to the 16-byte boundary in order to comply with the GPU memory alignment rules. Data blocks to be transferred should satisfy HLSL packing requirements in order to avoid memory corruption. For example, the **MeshMetadata** struct is explicitly padded in C# to guarantee exact 16-byte size:

```

1 [StructLayout(LayoutKind.Sequential)]
2 private struct MeshMetadata {
3     public uint vertexOffset; // 4 bytes
4     public uint indexOffset;  // 4 bytes
5     public uint hasGeometry;   // 4 bytes
6     public uint padding;       // 4 bytes (to force 16-byte alignment)
7 }

```

Moreover, GPU is expecting `float4` structures for efficient SIMD processing. Thus, the `_NormalMap` texture is using all four channels, where x, y, z correspond to the geometric coordinates while the alpha channel w acts as the binary coverage flag of valid geometry and padding.

The `SimulationMaterial` variables (reflectance, transmittance) are packed into the `Vector4` array and stored in the `_SimulationMaterials ComputeBuffer`. During the `IrradianceBaker` TLAS building using `rtas.AddInstance()`, special flags are set. Submeshes receive the `RayTracingSubMeshFlags.ClosestHitOnly` which is hardware-specific optimization disabling the use of Any-Hit shaders during the traversal, which increases silicon performance because the simulation requires only the nearest geometric intersection point.




Figure 5.4: HLSL Compute Shader modular structure and inclusion tree.

5.4.2. Texture Space Dispatch Paradigm

ChronoLux employs the Texture Space Ray Tracing paradigm [14], different from the classical screen-space approach. Instead of launching primary rays from the user's camera through the 1080p grid of the screen, the compute shader dispatches threads directly to the 2D (u, v) coordinates of the UV map of the target mesh (see Listing 9.1).

```

Require: PosMap, NormMap, SunDir,  $N_{samples}$ ,  $\Delta t$ 
for each thread id  $\in$  DispatchThreadID do
   $\mathbf{p} \leftarrow PosMap[id].xyz$ 
   $\mathbf{n} \leftarrow NormMap[id].xyz$ 
  if length( $\mathbf{n}$ ) < 0.1 then
    return // Terminate thread (empty UV padding)
  end if
   $E_{dir} \leftarrow 0$ 
   $NdotL \leftarrow \mathbf{n} \cdot SunDir$ 
  if  $NdotL > 0.0$  then
     $E_{dir} \leftarrow TRACEPATH(\mathbf{p}, SunDir, \dots, isDirectSun=true) \times NdotL$ 
  end if
   $E_{ind} \leftarrow 0$ 
  for  $i = 1$  to  $N_{samples}$  do
     $\omega_i \leftarrow GETCOSINEWEIGHTEDDIRECTION(\mathbf{n})$ 
     $E_{ind} \leftarrow E_{ind} + TRACEPATH(\mathbf{p}, \omega_i, \dots, isDirectSun=false)$ 
  end for
   $E_{total} \leftarrow E_{dir} + (E_{ind}/N_{samples}) \cdot \pi$ 
   $DoseMap[id] \leftarrow DoseMap[id] + E_{total} \cdot \Delta t$ 
   $DirectDoseMap[id] \leftarrow DirectDoseMap[id] + E_{dir} \cdot \Delta t$ 
   $IndirectDoseMap[id] \leftarrow IndirectDoseMap[id] + (E_{ind}/N_{samples}) \cdot \pi \cdot \Delta t$ 
end for

```

Figure 5.5: Pseudocode for Texture Space Ray Tracing dispatch separated into Direct and Stochastic Multi-Sampling paths.

The pre-pass bake stage operates via the standard rasterization pipeline to store world space positions (P) and geometric normals (N) into the `_PositionMap` and `_NormalMap` textures. As described in Algorithm 5.5 and detailed in Appendix 9.1, the ray-generation kernel (`IrradianceBakeRayGen`) is launched on the two-dimensional grid. If the calculated normal vector for a pixel has negligible length, this means there is a void area in the UV padding, and the thread ends prematurely to save computation cycles. Otherwise, the kernel uses P and N as the origin and direction to calculate the rendering equation. The architectural design of this solution allows separating light transport calculations from screen space while baking energy information into the architectural surface.

5.4.3. Pseudo-random Sequence Generation and Hashing

The principle behind Monte Carlo integration lies in the uniform random sampling over the integration domain (celestial hemisphere). In case of any clustering and correlations of the generated numbers, the calculation of the rendering equation becomes biased resulting in the appearance of non-uniform noise.

To handle this issue without managing large state arrays in the HLSL kernel, *ChronoLux* uses a stateless hashing scheme bound to the thread's pixel coordinates (x, y) and the simulation time step (Δt). The base seed is calculated as:

$$\text{baseSeed} = \text{Hash}(id.y \times \text{dims}.x + id.x + \text{_FrameIndex}) \quad (5.12)$$

By including the `_FrameIndex` (a monotonically increasing integer specifying the current time step) into the integer hash calculation, the generated pseudo-random multi-vector sequence gets constantly shifted within the duration of the simulation. As hundreds of time steps accumulate within the same `_DoseMap` texture buffer, the high-frequency noise cancels out mathematically, allowing natural convergence of the Monte Carlo estimation towards the true value of irradiance integral.

5.4.4. Direct Sunlight and Stochastic Multi-Sample Paths

To evaluate ambient light transport within the scene, *ChronoLux* splits the integration into two paths within the ray generation shader.

The first path calculates **Direct Sunlight**. If the surface normal matches the solar vector (`nDotSun > 0.0`), a single ray is shot directly towards the sun via `TracePath(isDirectSun = true)` invocation. Within `PathTracer.hlsl` file, direct rays are allowed to traverse transparent geometry. On material hit, the hardware queries the `_SimulationMaterials` buffer for the transmittance of the material. In case of non-zero transmittance (for example, stained glass) the throughput is scaled down according to it, the ray is slightly moved forward and the ray shoots further through the object towards the sun. If transmittance equals zero, the ray becomes an opaque shadow-caster and returns 0.0 lux.

The second path evaluates **Indirect Ambient Bouncing** via Stochastic Multi-Sampling. For each pixel, a `for`-loop generates $N_{samples}$ pseudo-random vectors distributed over the hemisphere using Cosine-Weighted Importance Sampling (`GetCosineWeightedDirection`) [13]. These rays call `TracePath(isDirectSun = false)`.

Unlike direct rays, indirect rays use the probabilistic **Russian Roulette** technique for path termination [6]. Russian Roulette offers an unbiased way of terminating the infinite path recursion without arbitrary light energy truncation (Listing 9.2). At each bounce intersection, a random float number $rv \in [0, 1)$ is generated from the hash sequence. Depending on the material's scalar reflectance (R) and transmittance (T) values, the ray's life span is decided according to the following cumulative threshold criteria:

1. **Reflection** ($rv < R$): The ray survives and spawns a new cosine-weighted hemisphere direction.
2. **Transmission** ($rv < R + T$): The ray passes through the material without changing the direction.
3. **Absorption** ($rv \geq R + T$): The ray dies probabilistically.

As the ray survives only if it passes the materials' characteristics check, the throughput is automatically used as the energy attenuation weighting factor, eliminating the necessity to manually scale the ray down according to the reflectance multiplier at each vertex. If an indirect ray gets out of the BVH without being absorbed, the `EvaluateSky()` function is called to receive the luminance of the atmosphere via the Perez Sky Model. Such highly recursive procedure allows *ChronoLux* to evaluate extremely occluded indirect lighting scenarios, for example, the light propagation through a corridor or the reflection from highly reflective marble floor onto textile artifact.

5.5. Advanced Mathematical Models and Optimization

In order to combine both scientific validity and computational feasibility, several advanced mathematical models and optimizations are implemented in the rendering en-

gine.

5.5.1. Perez Sky Model Atmospheric Coefficients

While the solar disk itself is responsible for the most intense irradiance contribution, a great part of the ambient lighting in the architecture is due to diffuse sky scattering. *ChronoLux* evaluates the latter using the Perez All-Weather Sky Model (see `SkyModel.hls1` in Listing 9.6), which describes highly anisotropic luminance of the celestial dome.

The implementation uses the clear sky scenario with Turbidity $T = 2.0$ and five empirically obtained static coefficients:

- $a = -0.1058$: Controls horizon brightening/darkening.
- $b = -0.0738$: Determines luminance gradient close to the horizon.
- $c = 1.8200$: Controls relative intensity of circumsolar corona.
- $d = -2.9744$: Determines spatial width of the circumsolar region.
- $e = 0.1778$: Controls backward-scattered component opposite to the sun.

When an indirect stochastic ray does not hit any geometry, the zenith angle θ and the angular separation γ from the sun are calculated and the relative luminance is evaluated using the scattering indicatrix function:

$$f(\theta, \gamma) = \left[1 + a \exp\left(\frac{b}{\cos \theta}\right) \right] \times [1 + c \exp(d\gamma) + e \cos^2 \gamma] \quad (5.13)$$

Finally, the result is normalized by the dynamic `DiffuseLux` value supplied from the CPU-side astrometry module via the normalization factor $E_{diffuse} \times 0.2$. Thus, the physically correct spectral distribution is obtained.

5.5.2. Hardware Traversal Mechanics (DXR and Face Culling)

Using the rendering engine based on the DirectX 12 DXR technology allows shifting the traversal bottleneck from software to hardware, leveraging the RT cores for the BVH traversal and Moeller-Trumbore ray-triangle intersection [11].

Also, when the `TraceRay` intrinsic is used for direct sunlight shadows calculation, a special ray flag is included in the payload: `RAY_FLAG_CULL_BACK_FACING_TRIANGLES`. As the direct shadow rays require only binary occlusion (visibility of the sun), and the architectural CAD models are usually closed manifolds, the flag quickly discards triangles pointing away from the ray origin. It reduces the computational load on RT cores by two times for shadow calculation. Please note, however, that this optimization is deliberately not used for indirect rays to allow the light reflection off the interior back-faces of transparent geometries, such as display cases.

5.6. Asynchronous Data Telemetry and UI Integration

5.6.1. Automated Sensor Placement

In physical metrology, lux meters are placed in the laboratory. Similarly, *ChronoLux* uses the `VirtualLuxSensor` objects. To perform objective baseline measurement, the `RuntimeModelLoader` automatizes the outdoor reference sensors' placement. Upon environment model loading, *ChronoLux* calls `CreateOutdoorReferenceSensors`. The latter

iterates over all renderers to create the global bounding box of the heritage site, calculating the maximum height Y_{max} , added 2.0m to it and the total geometric radius. Five sensors are created automatically – one in the center and four in the periphery – to guarantee that, regardless of CAD artifact complexity, unobstructed sensors will constantly measure the baseline sunlight potential for a given latitude and day.

5.6.2. Metrological Analytics Pipeline

Scientific metrology requires raw and extractable numeric data. To perform statistical analysis in Excel or Python, the precise light-dose metrics are required. The application uses the asynchronous telemetry pipeline to transfer high-resolution floating point data from VRAM to the CPU without affecting the 60 Hz rendering loop.

```
Require: DownsampleRT, DirectRT, IndirectRT
completedRequests  $\leftarrow$  0
Trigger AsyncGPUReadback.Request(DownsampleRT)
Trigger AsyncGPUReadback.Request(DirectRT)
Trigger AsyncGPUReadback.Request(IndirectRT)
while completedRequests < 3 do
    Wait asynchronously (Non-blocking)
end while
data  $\leftarrow$  Extract float arrays
MaxDose  $\leftarrow$  0,  $\mu$   $\leftarrow$  0, CoverageCount  $\leftarrow$  0,  $N_{active}$   $\leftarrow$  0
for  $i = 0$  to 65536 do    // Iterating 256x256 buffer
    val  $\leftarrow$  data[i]
    if val > 0.0 then
         $N_{active} \leftarrow N_{active} + 1$ 
         $\mu \leftarrow \mu + val$ 
        if val > MaxDose then MaxDose  $\leftarrow$  val
        end if
        if val > Threshold then CoverageCount  $\leftarrow$  CoverageCount + 1
        end if
    end if
end for
 $\mu \leftarrow \mu / N_{active}$ 
CoveragePercent  $\leftarrow$  (CoverageCount /  $N_{active}$ )  $\times$  100
```

Figure 5.6: CPU-Side Asynchronous Readback loop computing total scene variance and coverage metrics without blocking the primary rendering thread.

Iteration in sequence through a 1024×1024 float array on CPU at every simulated hour would considerably affect the performance. In order to solve this problem, `LightDoseSimulator.cs` leverages hardware-accelerated downsampling via `Graphics.Blit`. Thus, there are three significantly downscaled 256×256 RenderTextures (`_downsampleRT`, `_directDownsampleRT`, `_indirectDownsampleRT`) created on the GPU.

Instead of locking up the thread, the C# layer issues Unity's `AsyncGPUReadback.Request`. According to Algorithm 5.6, it launches the direct memory access (DMA) transfer. A closure callback keeps the `completed` counter. It is only after reaching value

3 that CPU fetches the underlying float arrays via `req.GetData<float>().ToArray()` which means all three buffers were transferred to the system RAM.

After transferring of float arrays to CPU memory, an optimized loop performs the processing of 65,536 downsampled pixels. It calculates the absolute maximum dose exposure, counts the strictly active pixels for deriving the arithmetic mean (μ), and determines precisely how many pixels are larger than the user-defined `coverageThreshold` (for instance, 50,000 Lux-Hours for delicate watercolor paintings). With the help of two different accumulation arrays, the telemetry framework allows isolating the exact sources of UV/Visible energy damage.

5.6.3. User Interface Toolkit and JSON Serialization Limitations

The graphical interface of *ChronoLux* is created using Unity’s contemporary UI Toolkit (UXML/USS), not the legacy Canvas immediate-mode GUI, thus employing the strict Model-View-Controller (MVC) DOM-based architecture. The responsive framework guarantees the proper scaling on 4K laboratory displays.

Scientific reproducibility implies that all the parameters of the experiment should be exactly savable. The `ProjectManager.cs` system uses the custom JSON serialization engine via `JsonUtility`. However, there is a noticeable software engineering limitation: Unity’s `JsonUtility` lacks serialization of standard C# `Dictionary<TKey, TValue>` collections.

In order to overcome this issue, without introducing the heavyweight third-party JSON libraries (thus increasing the size of the executable) the `ChronoProject` architecture uses a data flattening approach. Instead of serializing the dictionary mapping the sub-mesh’s string key to the material GUID, data is manually split into two perfectly synchronized parallel lists:

```
1 public List<string> objectNames = new List<string>();
2 public List<string> materialPresetNames = new List<string>();
```

When saving the experiment, the `ProjectManager` enumerates through the active simulation hierarchy, takes `RootName/MeshName` keys, appends them to `objectNames` list, and appends the corresponding Material GUID to the `materialPresetNames` list at exactly the same indices. While loading, the reconstruction loop builds the C# dictionary in the local memory and applies it to the DXR `IrradianceBaker`. The nested string-based hierarchy paths and material GUIDs make the JSON format human-readable experimental manifest.

Experiment launch implies the unidirectional data flow: the `LightDoseSimulator` calculates astrometric parameters, updates the internal C# model and issues UI property-changed events. The UI Controller catches these changes in order to update the DOM view. This design guarantees the separation of Model and View, thus avoiding direct mutation of the UI layout by intensive background computations.

5.7. Deep Engine Optimizations and Memory Integrity

Scalability of the dosimetric ray tracer to the level of millions of rays per frame requires the extensive optimization. The several low-level engine mechanisms need to be accurately synchronized in order to guarantee fidelity and stability of the hardware over the 8,760-hour simulation.

5.7.1. Ray-Origin Epsilon Offsetting and Floating Point Precision Limitations

There is a fundamental vulnerability in mathematics of the ray tracing: IEEE 754 floating point truncation. When the intersection is found by the BVH, the origin of the new ray is placed exactly on the geometric surface of the triangle. Due to the limitations of the 32-bit float precision, the calculated point can be marginally inside the geometry. Thus, the next bounce starting from this point can immediately re-intersect the origin triangle, causing noticeable shadow acne artifacts that destroy the dosimetric fidelity.

In order to avoid this, *ChronoLux* implements the strict normal-based epsilon offsetting in `PathTracer.hlsl`:

$$\mathbf{p}_{offset} = \mathbf{p}_{hit} + \mathbf{n}_{hit} \times 0.001 \quad (5.14)$$

It offsets the origin exactly for 1 millimeter (0.001 units in the normalized 2.0-meter space) along the geometric surface normal. Thus, the ray will be displaced outside of the surface before the next BVH traversal. Also, for the direct rays passing through the transparent material, the enlarged offset (`RAY_BIAS` $\times$ 5.0) is used in order to guarantee the passage through the thin window panes without internal trapping at the grazing angles.

5.7.2. Helmholtz Reciprocity and Energy Conservation Enforcement

Monte Carlo integration accuracy depends on the physical correctness of scene materials. Rendering Equation requires that the opaque surface cannot reflect more energy than it receives. For translucent materials (for instance, stained glass, thin Japanese paper), the sum of Reflectance (R) and Transmittance (T) should be less than 1.0. Otherwise, the system would produce energy out of nowhere, causing the runaway feedback and incorrect lux-hour accumulation.

ChronoLux imposes this law of physics at the C# binding layer. In `Irradiance-Baker.cs`, the `ClampEnergy` method works through the mathematical normalization but not through the aggressive hard clamping:

```

1 private static void ClampEnergy(ref float r, ref float t) {
2     r = Mathf.Clamp01(r);
3     t = Mathf.Clamp01(t);
4     float s = r + t;
5     if (s > 1f) {
6         r /= s;
7         t /= s;
8     }
9 }

```

It does not forcefully set the Transmittance to zero, but scales proportionally to keep the energy inside the 1.0 boundaries. These bounded values are passed to the `_SimulationMaterials` buffer, guaranteeing the throughput multiplier decay according to the universal conservation of energy.

5.7.3. Thread Synchronization and CPU-GPU Data Dependencies

ChronoLux heavily depends on the strict synchronization of the host CPU and the device GPU. Since the C# Coroutine loop is run asynchronously regarding the GPU command queue, race conditions become a major problem. If the CPU increments the

clock Δt and uploads new **SunDirection** uniforms while the GPU is running the prior hour's ray-tracing kernel, the light integration would become inaccurate.

In order to mitigate this, *ChronoLux* employs the disciplined pipeline with the help of Unity's **CommandBuffer** API. Instead of issuing the dangerous fire-and-forget **ComputeShader.Dispatch()**, **IrradianceBaker** queues all the **SetRayTracingVectorParam** and **DispatchRays** commands in the dedicated **CommandBuffer**. Synchronous submission via **Graphics.ExecuteCommandBuffer(cmd)** guarantees the accurate temporal ordering, meaning that the DXR kernels will use the correct sun position and solar intensity for the particular hour before advancing the master calendar loop of **Coroutine**.

5.7.4. VRAM Memory Integrity and Cleaning Protocol

Since the frequent swapping of the high-resolution 3D projects consumes the considerable amount of VRAM because of the **_doseMap** buffers, **_rtas** acceleration structures and the large **_vertexBuffer** arrays, the application is prone to the out-of-memory crashes if these resources are not explicitly freed.

In order to avoid such memory leaks, the explicit **CleanUp()** protocol is implemented in **IrradianceBaker**. Every time a new simulation starts or the **MonoBehaviour** gets destroyed, **Release()** is called for every **ComputeBuffer** and **RenderTexture**. It guarantees the stable laboratory operation without affecting the host system performance.

5.7.5. Data Export Payload Calculation and Mathematical Variance

At the end of every simulated chronological day, the high dynamic range accumulation buffer gets exported to the **.exr** file. It is achieved via reading the active dose map pixels to the **Texture2D** buffer via **ReadPixels**, conversion to EXR via **ImageConversion.EncodeToEXR(tex, Texture2D.EXRFlags.OutputAsFloat)**, and the synchronous **File.WriteAllBytes** to the local storage.

Also, the telemetry cache produced by the **TakeAsyncSnapshot** pipeline gets flushed to the CSV file via **ExportSimulationDataCSV** coroutine. The engine calculates the depths for the each saved frame, including the hourly variance (σ^2) across the surface of the artifact without risking the arithmetic overflow from millions of floating point sums. It is calculated as:

$$\sigma^2 = \left(\frac{\sum (\Delta_{dose})^2}{Area} \right) - \mu^2 \quad (5.15)$$

where variance is taken as the difference between the sum of squares and the square of the mean (μ). Thus, the metrics including the exact Delta Dose contributions (direct beam vs. indirect bouncing) are written to the **StreamWriter** and appended to **SimulationMetrics.csv**. This structure allows the external scientists to process the damage thresholds via any data science framework out of the Unity engine.

Chapter 6. Testing and Validation

6.1. Introduction to the Metrological Validation Framework

The transition from the pure computer graphics application to the scientifically justified Digital Twin requires a meticulous and exhaustive testing and validation process. Regarding the particular case of the *ChronoLux* engine, the validation should go beyond the subjective evaluation of the visual quality of the image and frame rates. The validation will focus on the integrity of radiometric light transport equations, the correctness of the calculations of solar ephemeris, and the stability of the Monte Carlo path tracing algorithm in extended 8,760-hour cycles.

This chapter offers a detailed empirical validation of the *ChronoLux* framework by means of a series of precisely defined cases. The goals are twofold: to illustrate the practical use of the software in the conservation scenarios; to prove mathematically the stability of the DirectX 12 compute shaders and C# orchestration in terms of the dosimetry data consistency.

The validation will be performed across four distinct conservation scenarios (A to D). They will allow isolating specific variables in the rendering equation like the geometric occlusion, the transparency, and albedo of material. In addition, there will be a special scientific validation case, which will monitor the spatial variance and the standard error of the mean (SEM) of the Monte Carlo integration.

6.2. Experimental Setup and Simulation Parameters

6.2.1. The Virtual Environment and Target Artifact

For empirical testing of the ability of the engine to quantify the cumulative light stress, a very detailed 3D geometric mesh is required. Stanford Bunny has been selected as the target artifact of these simulations. It has been re-topped and UV-mapped with a high-resolution 2K texture coordinate space (2048×2048). It is crucial due to the test of the Texture Space Ray Tracing paradigm, which forces the engine to launch more than 4.1 million (one per each UV texel) compute threads every frame.

The virtual museum has been built as a standard rectangular architectural space with a large window on the southern wall for sufficient solar penetration during daytime. The artifact is located on a pedestal inside this room.

6.2.2. Meteorological and Astrometric Configuration

The experiment has been performed in the setting of the execution of one full meteorological year of 8,760 solar hours. The geographical parameters were set to Latitude

44.427° N and Longitude 26.103° E, Bucharest, Romania. The selection of these coordinates allows getting a wide variation of the solar altitude in order to get maximum solar penetration in winter and minimum depth in summer.

Solar coordinates (Altitude and Azimuth) for each of the 8,760 hours have been calculated dynamically by the CPU using the `SunCalculator.cs` module. It is important to note that the mathematical correctness of these computations is crucial for the rendering equation because the direct illumination component depends on the accurate solar incident angle. The fundamental logic of the calculations lies in finding the solar declination (δ) and hour angle (ω) according to the Julian Day.

```

1 double declination = 0.409 * Math.Sin((2.0 * Math.PI / 365.0) * (
    julianDay - 1.39));
2 double hourAngle = (Math.PI / 12.0) * (solarTime - 12.0);
3
4 double elevationRad = Math.Asin(Math.Sin(latitudeRad) * Math.Sin(
    declination) +
5                               Math.Cos(latitudeRad) * Math.Cos(
    declination) * Math.Cos(hourAngle));
6 double elevationDeg = elevationRad * (180.0 / Math.PI);

```

Listing 6.1: Astrometric solar elevation calculation executed by the CPU orchestrator.

As you see, the obtained `elevationDeg` and corresponding Azimuth are transferred to the DirectX 12 compute shader via the constant buffer every frame in order to define the origin vector for the shadow rays.

6.2.3. Atmospheric Modeling via the Perez Model

Ambient skylight, which defines the indirect dose on the artifact, has been modeled by the Perez All-Weather Sky Model. In this stress testing validation, the atmosphere has been kept static with the clearness index (T) of 2.0, which means the completely clear sky. Although the real-life weather can include overcast, thus reducing irradiance, keeping it always clear constitutes an absolute worst case scenario. If the conservation strategy works in a year of constantly clear skies, it guarantees its effectiveness under the real life varying weather conditions.

6.2.4. Hardware and Execution Profiling

Simulations have been performed using a powerful PC with the NVIDIA GeForce RTX 4070 Ti Super GPU, AMD Ryzen 7 7800X3D CPU, and 32 GB of DDR5 6000 MT/s RAM. The hardware-accelerated DXR pipeline has been configured to emit 64 Monte Carlo rays per each UV texel per hourly frame. Despite the huge amount of computations (268 million rays per frame), the completely decoupled Texture Space architecture keeps a stable 60 FPS viewport, while the background C# telemetry thread asynchronously extracts floating-point accumulation buffers to the `SimulationMetrics.csv` file.

6.3. Data Telemetry and the Simulation Output

6.3.1. Downsampled Dosemap and GPU Readback

The extraction of radiometric data for each hour in the simulation of 8,760 hours poses a notable computational challenge. Extracting raw 2K resolution (2048×2048)

floating-point accumulation buffers directly from VRAM to the CPU side will saturate the PCI express bus and will heavily reduce the framerate. Therefore, in order to optimize the process, the *ChronoLux* engine employs the downsample architecture. The high-resolution `_DirectDoseMap` and `_IndirectDoseMap` buffers have been blitted into lower-resolution `_downsampleRT` render textures. Then, the `AsyncGPUReadback` API is used only on the downsampled maps to get a statistical representative yet light matrix of the dose distribution to the CPU telemetry thread without stopping the main 60 FPS rendering loop.

6.3.2. The Telemetry CSV Format

The *ChronoLux* engine progresses through the year and logs downsampled buffer data to the comma-separated values (CSV) log file every simulation tick. This CSV will be the dataset for all further empirical analysis and graph generation.

Table 6.1 illustrates sample telemetry data logged during Scenario A, alongside the mathematical derivation of each field.

Table 6.1: Sample telemetry data logged continuously during the simulation.

Time	Altitude	Azimuth	DeltaMax	DeltaAvg	MaxDose	AvgDose	Variance
Jan 01 09:27	4.31°	128.50°	151.15	42.56	151.15	42.56	900.68
Jan 01 09:57	8.33°	134.14°	276.93	69.90	409.00	112.46	2448.91
Jan 01 10:27	11.98°	140.10°	344.27	88.49	683.03	200.96	3911.17
Jan 01 10:57	15.18°	146.38°	398.31	101.83	1010.57	302.80	5208.42

The metrics being logged in the CSV are defined as follows:

- **Altitude / Azimuth:** the astrometric coordinates of the solar disc calculated for the current Julian Day.
- **DeltaMax / DeltaAvg:** the light energy (in Lux-Hours) accumulated *only* during the current time step, isolating the instantaneous hourly stress.
- **MaxDose / AvgDose:** the running cumulative total of light energy throughout the whole simulation history. **MaxDose** refers to the single highest texel (hotspot), while **AvgDose** refers to the mean value on the valid surface area.
- **Variance (σ^2):** the spatial variance of the downsampled light distribution to mathematically prove the physical difference between shadowed and exposed areas as the light accumulates.

Relying on the mentioned exhaustive and continuous logging system rather than on the static snapshots allows conservators to trace the exact moments of the peak stress and correlate them with seasonal and diurnal changes.

6.4. Overview of Conservation Scenarios

Before proceeding to the empirical findings, four distinct conservation scenarios (A to D) will be described, constituting the core of this validation. The gradual introduction of different geometric and material interventions will be tested here. Scenario A is the baseline untreated museum environment with the artifact fully exposed to the clear southern window and highly reflective architectural materials. Scenario B focuses on the material filtration by introducing the UV-blocking conservation glass and matte flooring while keeping the highly exposed position of the artifact. Scenario C does the opposite

thing: it examines purely geometric occlusion, moving the artifact to the deeply shaded corner while reverting materials to the highly reflective baseline. Scenario D combines the corner occlusion and the protection of tinted glass and matte flooring. Further sections will analyze the telemetry data and the physical light transport mechanisms of each scenario separately.

6.5. Scenario A: The Baseline Exposure

6.5.1. Configuration and Untreated Conditions

Scenario A is the experimental control, which defines the baseline light exposure of the artifact in the untreated museum environment. The goal is to quantify the photodegradation potential of this exposure without the conservation interventions.

Architectural materials inside the UXML interface have been configured to represent highly reflective raw building materials. The floor had the high Lambertian albedo of $R = 0.50$, meaning the polished light-colored surface, which maximizes the indirect light diffusion. The southern window has been modeled with a standard clear glass material (Transmittance $T = 0.84$ and Reflectance $R = 0.08$), letting 84% of direct solar energy pass without attenuation. The artifact has been positioned in the path of the window for maximum exposure.

6.5.2. Empirical Analysis of the Catastrophic Dose

The simulations have been started in *ChronoLux* engine, accumulating 8,760 hours of the 2026 meteorological year. Extraction of telemetry data from the `_DirectDoseMap` and `_IndirectDoseMap` GPU buffers has shown a catastrophic light stress.

Figures 6.1 and 6.2 illustrating cumulative dose metrics show a peak cumulative light dose of 6.6×10^6 Lux-Hours, largely exceeding the CIE upper recommendation of 50,000 Lux-Hours per year for highly sensitive materials. The average cumulative dose across the entire surface of the artifact has also demonstrated unsustainable exponential growth, which confirms the existence of the photodegradation risk across the asset.

6.5.3. Diurnal Fluctuation and Hourly Stress

In order to explain the mechanisms of this energy accumulation, diurnal and seasonal irradiance fluctuations have been analyzed using the diurnal heatmap.

Figure 6.3 represents the heatmap of hourly exposure in a 365-day cycle, showcasing the high-frequency peaks near the solar noon and seasonal elevation changes.

It is governed by the solar elevation: in winter, the sun goes much lower on the horizon, which allows the direct sunlight vectors to penetrate deeper via the southern window and hit the artifact. In summer, the high position of the sun leads to the indirect dosing, which dominates due to the floor reflection, but the winter dominance is evident in total accumulation.




Figure 6.1: CIE Light Dosimetry results for Scenario A reveal cumulative light stress exposure over 8760 hours resulting in a large peak light dose of 6.6M Lux-Hours far greater than the maximum recommended light stress of 50,000 Lux-Hours.




Figure 6.2: CIE Light Dosimetry results for Scenario A showing cumulative average dose over the 8760 hours. The exponential growth curve illustrates the unmitigated accumulation of energy.




Figure 6.3: A heat map of diurnal exposure illustrates hourly exposure levels in a 365-day cycle, displaying high frequency peaks corresponding to solar noon hours and seasonal solar elevation changes.

6.6. Scenario B: Material Intervention

6.6.1. Simulating Physical Filtration

Expanding upon Scenario A, Scenario B tests the influence of the material interventions while keeping the position of the artifact the same. The clear window glass is being replaced with the UV-blocking conservation tinted glass (Transmittance $T = 0.45$ and Reflectance $R = 0.12$), while the floor material is changed to the dark matte with the albedo $R = 0.15$.

6.6.2. The Physics of Ray Attenuation

The mentioned change of materials substantially affects the Monte Carlo path tracing in the DirectX 12 compute shader. Shadow rays toward the solar disc to calculate the direct dose E_{dir} traverse the BVH and intersect the window geometry. Since the window is transmissive, the ray energy is multiplied by the transmittance (0.45), which reduces the direct solar energy by 55% before reaching the artifact.

The NEE rays into the Cosine Hemisphere for the indirect dose E_{ind} lose energy while interacting with the matte floor, whose low albedo acts as the energy sink, preventing light from entering back to the lower surfaces of the artifact.

6.6.3. Results of the Material Intervention

Scenario B shows the drastic decrease in the light stress, with the peak dose falling from 6.6 millions to about 150 thousands Lux-Hours (97.7% reduction). Nevertheless, it is still above the CIE 50,000 Lux-Hour limit for highly sensitive artifacts, which means that material filtration alone is not enough in case of the exposed position of the artifact to the southern window.

6.7. Scenario C: Positional Optimization and Geometric Occlusion

6.7.1. The Role of Architectural Shadows

Scenario C evaluates the influence of the spatial position and geometric occlusion. Materials have been reverted to the highly reflective baselines ($T = 0.84$ and $R = 0.50$), while the artifact is moved to the deeply shaded corner, distant from the direct southern window illumination.

6.7.2. Phenomenological Analysis: Window vs Corner Geometry

The phenomenological comparison of the hourly peak doses has been made for a 10-day window around the winter solstice.




Figure 6.4: The 10-day winter solstice comparison of peak hourly doses for artifacts placed next to windows versus artifacts placed deep in room corners. The window placement is dominated by direct solar energy, while the corner placement is exclusively driven by diffuse ambient skylight.

As shown in Figure 6.4, the window-adjacent artifact has direct solar driven hourly dose spikes near 30,000 Lux-Hours at the solar noon, while the corner-located one has zero direct dose, with the 5,000 Lux-Hour dose being caused exclusively by the indirect diffuse skylight, which is reflected by the walls and floor. In the HLSL compute shader, the shadow ray to the sun intersecting a solid wall terminates ($t > 0$, $T = 0$), thus giving $E_{dir} = 0$. Thus, this validates the ability of the engine to distinguish direct and indirect transport via BVH intersections, avoiding the rasterization-based glowing artifacts.

6.8. Scenario D: Combined Optimal Conservation

6.8.1. Synergistic Interventions

Scenario D implements the optimal strategy: the geometric occlusion (corner placement) in combination with the material filtration (tinted glass $T = 0.45$ and matte floor $R = 0.15$).




Figure 6.5: Comparative study of peak light dose across the four scenarios. Scenario D demonstrates the synergistic effectiveness of combining geometric occlusion with protective material interventions, achieving a dose safely below the 50,000 Lux-Hour limit.

Figure 6.5 represents the bar plot comparison of the peak light doses across Scenarios A to D, demonstrating that Scenario D reaches the conservation objective: the total peak dose over 8,760 hours falls significantly below the 50,000 Lux-Hour threshold.

The synergy of effects demonstrates the value of the *ChronoLux* Digital Twin. The framework allows evaluating the combined interventions before any physical changes will be made, reducing risks and costs in conservation planning.

6.9. Scientific Validation: Variance and Convergence

6.9.1. The Necessity of Mathematical Integrity

Although the heatmaps and CSV telemetry provide actionable data, their scientific validity needs to be verified. The *ChronoLux* engine uses the Monte Carlo stochastic integration method, which inevitably introduces the variance (or noise) in the sampling. Ineffective sample numbers and the biased PRNGs may lead to the inaccurate dosimetry, which will decrease the metrological validity of the Digital Twin.

In order to verify the integration, the network of virtual Lux Sensors performs the asynchronous sampling of the raw 32-bit floating-point accumulation buffers to compute spatial variance and the convergence error during runtime.

6.9.2. Spatial Dose Variance Tracking

Spatial variance (σ^2) is the measure of the deviation of individual texel doses from the mean value. The engine computes it as follows:

$$\sigma^2 = \frac{1}{N} \sum_{i=1}^N (D_i - \bar{D})^2 \quad (6.1)$$

Where N is the number of valid surface texels, D_i is the dose at the texel i , and $\bar{D}$ is the mean dose.




Figure 6.6: Continuous monitoring of light spatial variance (σ^2) across the 8760-hour simulation. The variance directly tracks the physical accumulation of light stress on exposed geometry.

Figure 6.6 represents the tracking of spatial variance across the 8,760-hour simulation, showcasing the stable upward trend with the distinctive steps in winter months, which is consistent with the physical expectations, since the exposed regions accumulate energy, while occluded regions stay dark.

6.9.3. Monte Carlo Convergence Error (SEM)

In order to verify that 64 rays per texel are enough to make the integration stable, the standard error of the mean (SEM) is computed:

$$SEM = \frac{\sigma}{\sqrt{N_{rays}}} \quad (6.2)$$

Figure 6.7 demonstrates the gradual decrease and stabilization near zero of SEM, which means a healthy Monte Carlo estimator.




Figure 6.7: Virtual Sensor Monte Carlo convergence error (SEM) steadily dropping and stabilizing near zero, confirming the mathematical reliability of the GPU-based path tracing integration.

6.9.4. Asynchronous Telemetry Pipeline Integrity

In order to keep the performance of the engine, the telemetry pipeline is asynchronous. Locked GPU readback of a 2K buffer per frame degrades the performance; therefore, the *ChronoLux* engine uses `AsyncGPUReadback` API. The data requests are issued to the GPU, while the CPU thread continues its work, making possible the 60 FPS viewport rendering without any interruption.

```

1 IEnumerator ExtractTelemetryAsync() {
2     var request = AsyncGPUReadback.Request(accumulationBuffer);
3     while (!request.done) {
4         yield return null; // Yield to main thread, keep 60 FPS
5     }
6
7     if (request.hasError) {
8         Debug.LogError("VRAM Readback Failed.");
9         yield break;
10    }
11
12    var data = request.GetData<float>();
13    ProcessVarianceMetrics(data);
14 }

```

Listing 6.2: Asynchronous background thread in `LightDoseSimulator.cs` ensuring non-blocking data extraction.

As shown above, the CPU thread processes variance metrics only after the GPU tells that the transfer is completed, which ensures that the CSV files for validation graph

generation contain complete radiometric data independently of the rendering.

6.10. Conclusion of the Empirical Validation

The comprehensive empirical validation proves that the *ChronoLux* framework represents a scientifically sound instrument for the predictive light dosimetry. Correctly reproducing architectural occlusion physics, modeling energy attenuation through the refractive materials, and mathematically proving the convergence of the Monte Carlo core makes the engine to elevate from the graphics prototype to the metrological Digital Twin suitable for use in the contemporary cultural heritage conservation practices.

Chapter 7. User’s Manual

7.1. Introduction

The *ChronoLux* Digital Twin is designed as a standalone, high-performance desktop application. It integrates a real-time 3D viewport with a comprehensive analytical dashboard, empowering conservators and metrologists to simulate, visualize, and quantify cumulative light stress on cultural heritage assets. This chapter provides a detailed, step-by-step guide to installing the software, navigating the virtual laboratory, configuring simulation parameters, and exporting scientifically valid dosimetry data.

7.2. Installation and System Requirements

Because *ChronoLux* relies on hardware-accelerated DirectX 12 compute shaders for Texture Space Ray Tracing, it requires a modern workstation capable of handling high-density parallel workloads.

7.2.1. Hardware Prerequisites

- **GPU:** NVIDIA GeForce RTX 30-series or 40-series (RTX 4070 Ti Super or higher recommended) with full DXR (DirectX Raytracing) support.
- **CPU:** A multi-core processor (e.g., AMD Ryzen 7 7800X3D) to handle the asynchronous telemetry extraction without stalling the GPU.
- **RAM:** Minimum 16 GB, recommended 32 GB DDR5 for caching high-resolution floating-point render textures.
- **Storage:** SSD storage for fast retrieval of the 2K and 4K uncompressed texture atlases.

7.2.2. Software Installation and Artifact Import

The application is distributed as a pre-compiled Unity executable. No additional dependencies or external runtimes are required beyond standard Windows 10/11 drivers.

A critical feature of the *ChronoLux* framework is its robust asset pipeline. The environment is not restricted to a single monolithic mesh. The user can import complex environments comprising multiple disparate objects (e.g., distinct walls, windows, floors, pedestals, and target artifacts). This allows for granular material assignment and precise architectural modeling, enabling the software to calculate exact geometric occlusions and secondary light bounces across multiple surfaces independently.

7.3. The Project Launcher Dashboard

Upon launching the executable, the user is greeted by the Project Launcher Dashboard. This interface manages the multi-project ecosystem, ensuring that different museum wings, artifacts, or hypothetical conservation scenarios can be saved and loaded independently.




Figure 7.1: The Project Launcher Dashboard, allowing for the creation and management of independent simulation profiles via JSON persistence.

As shown in Figure 7.1, the launcher lists all previously created projects. Project states are persistently serialized into JSON files stored within the local application directory. This persistence guarantees that all material assignments, artifact transformations, and geographic coordinates are perfectly retained across sessions. To begin, the user may select an existing profile or instantiate a new simulation environment.

7.4. Viewport Navigation and Interaction

Entering a project loads the 3D Virtual Laboratory. The user interface within the laboratory is divided into the central 3D viewport and the surrounding analytical UXML panels.

Navigation is handled via the **FreeLookCamera** system. To look around the environment, the user must click and drag the mouse within the viewport area. Movement is achieved using the standard **W**, **A**, **S**, **D** keyboard configuration. For rapid traversal of large architectural spaces, holding the **Shift** key activates a 10x velocity multiplier.

At the center of the screen, a precise crosshair facilitates the **ObjectPicker** raycast system. Aligning the crosshair with any sub-mesh in the environment and clicking will select that specific object, opening its parameters in the right-hand Material Configuration panel. The cursor lock can be dynamically toggled at any time by pressing the **ESC** key, freeing the mouse to interact with the surrounding sliders and buttons.




Figure 7.2: The main 3D viewport showcasing the target artifact and the central crosshair used for the `ObjectPicker` interaction system.



Figure 7.3: The false-color "Inferno" heatmap visualization mapping peak light stress directly onto the artifact's 3D geometry in real-time.

7.5. The Laboratory HUD and Telemetry

The left side of the screen houses the live telemetry and transformation Heads-Up Display (HUD). This panel provides real-time feedback on both the spatial configuration of the selected artifact and the ongoing radiometric calculations.

As illustrated in Figure 7.4, the HUD contains several critical subsystems:

- **Transformation Controls:** Granular input fields allowing the user to precisely override the Position, Rotation, and Scale of the selected object.
- **Material Properties:** Direct readout of the selected object's assigned physical properties.
- **Astrometric Telemetry:** Live updates of the Solar Altitude and Azimuth angles based on the active Julian Day.
- **Dosimetry Metrics:** Continuous tracking of the Peak Dose (maximum hotspot accumulation) and the average dose across the mesh.
- **Scientific Validation:** A real-time readout of the virtual sensor error (variance), confirming the stability of the Monte Carlo integration.
- **Radiation Waveform Graph:** A dynamic, scrolling waveform visualization that plots the intensity of the light radiation per hour, instantly reflecting diurnal spikes during solar noon.
- **Data Export:** A dedicated button for triggering the asynchronous export of the simulation data to a CSV file.

7.6. Material Configuration and Catalog

Selecting any object in the viewport via the crosshair activates the Material Configuration catalog on the right-hand side of the interface. This catalog is crucial for setting up the preventive conservation scenarios (e.g., swapping clear glass for UV-tinted glass).

The interface (Figure 7.6) provides intuitive, live-updating sliders to adjust the Reflectance (R) and Transmittance (T) coefficients of the selected mesh. A strict energy conservation clamp ensures that $R + T \leq 1.0$, preventing the creation of physically




Figure 7.4: The Laboratory HUD displaying real-time transformation controls, astrometric telemetry, Monte Carlo variance tracking, and the dynamic hourly radiation waveform.



Figure 7.5: The Simulation Configuration panel managing the orchestration of the 8,760-hour loop, featuring execution controls, a progression bar, and a daily accumulation slider.




Figure 7.6: The Material Catalog interface, featuring real-time Reflectance and Transmittance sliders and the texture library.

impossible materials.

Additionally, the catalog features a built-in library of scientifically validated CC0 PBR (Physically Based Rendering) textures. The user can visually identify and apply high-fidelity textures such as Brick, Clear Glass, Mirror, Plaster, Hardwood, Carpet, and Grass by clicking the corresponding preview buttons. These textures are immediately mapped to the object's UV space, instantly updating its interaction with the Monte Carlo light rays.

7.7. Simulation Configuration and Execution

Once the environment geometry and materials are properly configured, the user must initialize the time-stepping simulation loop. This is managed via the Simulation Configuration panel.

The Simulation Configuration panel (Figure 7.5) provides complete control over the temporal execution of the Digital Twin:

- **Execution Controls:** Dedicated buttons to **Start**, **Bake**, or **Stop** the simulation. The **Clear Map** button flushes the GPU VRAM accumulation buffers, resetting the dose to zero.
- **Progression Bar:** A visual loading bar that tracks the progression of the simulation through the 8,760 hours of the meteorological year.
- **Daily Accumulation Slider:** An interactive time slider that allows the user to manually scrub through the simulation, visualizing the accumulation of dose on a per-day basis, providing a granular look at the temporal stress distribution.

7.8. Heatmap Visualization and Data Export

As the simulation progresses or after it successfully completes the full year cycle, the analytical data must be visualized and exported.

7.8.1. Viewport Heatmap

The engine features a sophisticated false-color heatmap shader, mapped directly onto the geometry of the target artifact (Figure 7.3). Using a high-contrast "Inferno" palette (transitioning from deep purple through red to bright yellow), it visually isolates the spatial hotspots receiving the highest dose. This allows conservators to intuitively understand which physical areas of the object are at maximum risk without having to read raw numeric data.

7.8.2. Headless Telemetry Export

For rigorous academic and metrological analysis, visual heatmaps are insufficient. By utilizing the export functionality, the asynchronous C# orchestration thread writes the entire year's worth of accumulated data to a robust CSV log file. This file contains the hourly breakdowns of Altitude, Azimuth, Peak Dose, Average Dose, and Spatial Variance. Simultaneously, high dynamic range EXR image snapshots of the accumulation buffers are exported. These standardized formats ensure that the raw scientific data generated by *ChronoLux* can be natively imported into external tools like Microsoft Excel, Python, or MATLAB for further graphing, comparative study, and publication formatting.

Chapter 8. Conclusions

This thesis aims to develop and deploy an accurate Heritage Digital Twin capable of simulating the cumulative effects of light stress on heritage objects with mathematical precision through *ChronoLux*, a scientifically-grounded metrological digital twin tool. The fundamental assumption of this tool is unifying the direct and indirect light calculations with the help of a Monte Carlo path tracer integrated with Next Event Estimation (NEE) and the Perez All-Weather Sky Model.

The main innovation in the system is a Texture Space Ray Tracing pipeline, which allows separating the dosimetry calculation process from the traditional screen-space ray tracing approach. Thanks to the Texture Space Ray Tracing, it is possible to perform the exact geometry-based occlusions and secondary bounce calculations regardless of the view position. The implementation utilizes the modern hardware-accelerated compute shader technology (DirectX 12), making it capable of dealing with a high amount of parallel tasks efficiently. This performance, along with the headless telemetry export functionality, ensures that the raw dose accumulations during an 8,760-hour meteorological year are converted into easily parseable CSV files for further research.

Tests conducted on the defined conservation scenarios confirmed the effectiveness of the tool and its stability in terms of performance and results. The first scenario was used as a baseline, where the simulation yielded 6.6 million Lux Hours accumulation, which was consistent with the theoretical expectations. In scenarios B, C, and D, the correctness of the material interaction and positional optimization systems was verified. Specifically, the combined interventions in Scenario D reduced the peak dose to 150,000 Lux Hours—a 97.7% reduction—while strictly observing the energy conservation boundaries ($R + T \leq 1.0$). It is worth noting that the continuous monitoring of the virtual sensor error indicated the stability of Monte Carlo variance throughout the entire process, which proves that the generated false-color "Inferno" heatmaps indeed pinpoint the spatial hotspots with the maximum dose exposure risk.

While the implemented version of *ChronoLux* is highly efficient, there are some limitations in the scope related to hardware and meteorological assumptions. Firstly, the utilization of high-end DXR-compatible hardware (such as NVIDIA RTX 30 or 40 series graphics cards) makes the access to the tool limited to institutions having their own workstations. Secondly, the present environmental simulations assume the absence of clouds, hence overestimating the annual light exposure for geographic regions having a lot of rainfalls or a cloudy climate. Lastly, the physical material interactions use the standard BRDF model and cannot account for the subsurface scattering processes occurring with complex translucent artifacts.

Overstepping these simulation boundaries is a crucial next step for the development of the Digital Twin. The main directions for the future development include incorporation of the cloud-cover data into the Perez Sky Model to make the annual dose estimates more precise and localized. Additionally, migration of the compute shaders to the cloud-

based infrastructure will allow to get rid of hardware dependencies and increase platform compatibility for conservators. Expansion of the material library to cover the subsurface scattering will be the last stage in developing the comprehensive preventative conservation simulator.

Bibliography

- [1] S. Michalski, “Damage to museum objects by visible radiation (light) and ultraviolet radiation (uv),” *Conference on lighting in museums, galleries and historic houses*, pp. 3–16, 1987. [Online]. Available: <https://www.canada.ca/en/conservation-institute/services/agents-deterioration/light.html>
- [2] CIE, “Control of damage to museum objects by optical radiation,” International Commission on Illumination, Tech. Rep. 157:2004, 2004. [Online]. Available: <https://cie.co.at/publications/control-damage-museum-objects-optical-radiation>
- [3] H. Entradas Silva, G. Coelho, and F. Henriques, “Climate monitoring in world heritage list buildings with low-cost data loggers: The case of the jerónimos monastery in lisbon (portugal),” *Journal of Building Engineering*, p. 101029, 10 2019. [Online]. Available: <https://doi.org/10.1016/j.job.2019.101029>
- [4] A. Chianese, F. Piccialli, and J. E. Jung, “The internet of cultural things: Towards a smart cultural heritage,” in *2016 12th International Conference on Signal-Image Technology & Internet-Based Systems (SITIS)*, 2016, pp. 493–496. [Online]. Available: <https://doi.org/10.1109/SITIS.2016.83>
- [5] J. T. Kajiya, “The rendering equation,” *SIGGRAPH Comput. Graph.*, vol. 20, no. 4, p. 143–150, Aug. 1986. [Online]. Available: <https://doi.org/10.1145/15886.15902>
- [6] M. Pharr, W. Jakob, and G. Humphreys, *Physically Based Rendering: From Theory to Implementation*. MIT Press, 2023. [Online]. Available: <https://books.google.ro/books?id=IQpmzwEACAAJ>
- [7] G. Ward Larson and R. Shakespeare, *Rendering with Radiance: the art and science of lighting visualization*. Morgan Kaufmann, 1998. [Online]. Available: <https://www.radiance-online.org/learning/documentation/book>
- [8] P.-P. Sloan, J. Kautz, and J. Snyder, “Precomputed radiance transfer for real-time rendering in dynamic, low-frequency lighting environments,” *ACM Trans. Graph.*, vol. 21, no. 3, pp. 527–536, Jul. 2002. [Online]. Available: <https://doi.org/10.1145/566654.566612>
- [9] P. Jouan and P. Hallot, “Digital twin: Research framework to support preventive conservation policies,” *ISPRS International Journal of Geo-Information*, vol. 9, no. 4, 2020. [Online]. Available: <https://www.mdpi.com/2220-9964/9/4/228>
- [10] F. Niccolucci, A. Felicetti, and S. Hermon, “Digital twins for cultural heritage,” *Heritage*, vol. 5, no. 4, pp. 2397–2414, 2022. [Online]. Available: <http://dx.doi.org/10.2139/ssrn.4274457>

-
- [11] J. Burgess, “Rtx on - the nvidia turing gpu,” *IEEE Micro*, vol. PP, pp. 1–1, 02 2020. [Online]. Available: https://www.researchgate.net/publication/339036275-RTX_On_-_The_NVIDIA_Turing_GPU
 - [12] R. Perez, R. Seals, and J. Michalsky, “All-weather model for sky luminance distribution—preliminary configuration and validation,” *Solar Energy*, vol. 50, no. 3, pp. 235–245, 1993. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/0038092X9390017I>
 - [13] E. Veach and L. J. Guibas, “Optimally combining sampling techniques for monte carlo rendering,” in *Proceedings of the 22nd Annual Conference on Computer Graphics and Interactive Techniques*, ser. SIGGRAPH ’95. New York, NY, USA: Association for Computing Machinery, 1995, pp. 419–428. [Online]. Available: <https://doi.org/10.1145/218380.218498>
 - [14] J. Munkberg, J. Hasselgren, P. Clarberg, M. Andersson, and T. Akenine-Möller, “Texture space caching and reconstruction for ray tracing,” *ACM Trans. Graph.*, vol. 35, no. 6, Dec. 2016. [Online]. Available: <https://doi.org/10.1145/2980179.2982407>
 - [15] I. Reda and A. Andreas, “Solar position algorithm for solar radiation applications,” *Solar Energy*, vol. 76, no. 5, pp. 577–589, 2004. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0038092X0300450X>
 - [16] M. Blanco-Muriel, D. C. Alarcón-Padilla, T. López-Moratalla, and M. Lara-Coira, “Computing the solar vector,” *Solar Energy*, vol. 70, no. 5, pp. 431–441, 2001. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0038092X00001560>
 - [17] J. Meeus, *Astronomical Algorithms*. Willmann-Bell, 1991. [Online]. Available: <https://books.google.ro/books?id=1smAQgAACAAJ>
 - [18] F. Kasten and A. T. Young, “Revised optical air mass tables and approximation formula,” *Appl. Opt.*, vol. 28, no. 22, pp. 4735–4738, Nov 1989. [Online]. Available: <https://opg.optica.org/ao/abstract.cfm?URI=ao-28-22-4735>
 - [19] M. Iqbal, *An introduction to solar radiation*. Academic Press, Orlando, FL, 01 1983. [Online]. Available: <https://www.osti.gov/biblio/5596615>

Chapter 9. Appendix

9.1. Core HLSL Path Tracing Implementation

The following code details the exact DirectX Raytracing implementation used within *ChronoLux* to evaluate the dosimetric simulation. The integration is explicitly split into direct sun evaluation and a secondary stochastic multi-sampling loop for indirect ambient bouncing.

```
1 [shader("raygeneration")]
2 void IrradianceBakeRayGen()
3 {
4     uint2 id = DispatchRaysIndex().xy;
5     uint2 dims = DispatchRaysDimensions().xy;
6     float4 bakedPos = _PositionMap[id];
7     float4 bakedNormal = _NormalMap[id];
8
9     if (bakedPos.w < 0.01 || length(bakedNormal.xyz) < 0.1) return;
10
11     float3 origin = bakedPos.xyz;
12     float3 normal = normalize(bakedNormal.xyz);
13
14     uint baseSeed = Hash(id.y * dims.x + id.x + _FrameIndex);
15     float totalIrradiance = 0.0;
16     float nDotSun = dot(normal, _SunDirection);
17
18     // 1. Direct Sun
19     float directIrradiance = 0.0;
20     if (nDotSun > 0.0)
21     {
22         uint sunSeed = baseSeed;
23         directIrradiance = TracePath(origin, _SunDirection, normal,
24         sunSeed, true, _SceneRTAS, _SunDirection, _BeamLux, _DiffuseLux) *
25         nDotSun;
26         totalIrradiance += directIrradiance;
27     }
28
29     // 2. Stochastic Multi-Sampling
30     float indirectIrradiance = 0.0;
31     uint numSamples = max(1, _SamplesPerPixel);
32     for (uint i = 0; i < numSamples; i++)
33     {
34         uint sampleSeed = Hash(baseSeed + i);
35         float3 randomDir = GetCosineWeightedDirection(normal,
36         sampleSeed, PI);
37         indirectIrradiance += TracePath(origin, randomDir, normal,
38         sampleSeed, false, _SceneRTAS, _SunDirection, _BeamLux, _DiffuseLux)
39         ;
40     }
```

```

36
37     float avgIndirect = (indirectIrradiance / (float)numSamples) * PI;
38     totalIrradiance += avgIndirect;
39
40     _DoseMap[id] += (totalIrradiance * _DeltaHours);
41     _DirectDoseMap[id] += (directIrradiance * _DeltaHours);
42     _IndirectDoseMap[id] += (avgIndirect * _DeltaHours);
43 }

```

Listing 9.1: IrradianceBake.raytrace: The primary ray generation kernel mapping Target Mesh UVs to stochastic integration over the hemisphere.

```

1 float TracePath(float3 origin, float3 direction, float3 surfaceNormal,
  inout uint seed, bool isDirectSun, RaytracingAccelerationStructure
  rtas, float3 sunDir, float beamLux, float diffuseLux)
2 {
3     float throughput = 1.0;
4     float3 currentOrigin = origin + surfaceNormal * RAY_BIAS;
5     float3 currentDir = direction;
6
7     [loop]
8     for (uint bounce = 0; bounce < MAX_BOUNCES; bounce++)
9     {
10         RayDesc ray;
11         ray.Origin = currentOrigin;
12         ray.Direction = currentDir;
13         ray.TMin = 0.001;
14         ray.TMax = 10000.0;
15
16         Payload payload;
17         payload.isOccluded = 1;
18         payload.reflectance = 0.0;
19         payload.transmittance = 0.0;
20         payload.hitDistance = 0.0;
21         payload.worldNormal = float3(0, 1, 0);
22
23         uint rayFlags = RAY_FLAG_FORCE_OPAQUE;
24         if (isDirectSun) rayFlags |=
RAY_FLAG_CULL_BACK_FACING_TRIANGLES;
25
26         TraceRay(rtas, rayFlags, 0xFF, 0, 1, 0, ray, payload);
27
28         if (payload.isOccluded == 0)
29         {
30             return EvaluateSky(currentDir, sunDir, beamLux, diffuseLux,
isDirectSun) * throughput;
31         }
32
33         if (isDirectSun)
34         {
35             float trans = saturate(payload.transmittance);
36             if (trans <= 0.0) return 0.0; // Blocked by opaque geometry
37
38             throughput *= trans;
39             // Use a larger bias to guarantee escaping the pane
geometry, especially at grazing angles
40             currentOrigin += currentDir * (payload.hitDistance +
RAY_BIAS * 5.0);

```

```

41         continue;
42     }
43
44     // Probabilistic event selection (Russian Roulette)
45     float rv = Random(seed);
46     float refl = saturate(payload.reflectance);
47     float trans = saturate(payload.transmittance);
48
49     if (rv < refl)
50     {
51         float3 hitPoint = currentOrigin + currentDir * payload.
hitDistance;
52         currentDir = GetCosineWeightedDirection(payload.worldNormal
, seed, PI);
53         currentOrigin = hitPoint + payload.worldNormal * RAY_BIAS;
54     }
55     else if (rv < refl + trans)
56     {
57         currentOrigin += currentDir * (payload.hitDistance +
RAY_BIAS);
58     }
59     else
60     {
61         return 0.0;
62     }
63
64     if (throughput <= MIN_VISIBILITY) break;
65 }
66 return 0.0;
67 }

```

Listing 9.2: PathTracer.hlsl: The core recursive path tracing loop evaluating Russian Roulette absorption and transparent geometric traversal.

9.2. C# Orchestrator Implementations

The following code snippets detail the core C# CPU orchestration logic. Listing 9.3 demonstrates the dynamic chronological time-stepping Coroutine used to drive the simulation safely across the OS thread schedules.

```

1 private IEnumerator RunSimulationInternal() {
2     float deltaHours = stepSeconds / 3600f;
3     int totalDays = endDay - startDay + 1;
4     int currentDayIdx = 0;
5
6     string exportDir = Path.Combine(Application.dataPath, "../Exports",
DateTime.Now.ToString("yyyyMMdd_HH:mm:ss"));
7     Directory.CreateDirectory(exportDir);
8
9     for (int day = startDay; day <= endDay; day++) {
10         DateTime date = new DateTime(year, 1, 1).AddDays(day - 1);
11         currentDayIdx++;
12
13         DateTime sunrise = SunCalculator.FindSunrise(date, latitude,
longitude, utcOffset);
14         DateTime sunset = SunCalculator.FindSunset(date, latitude,
longitude, utcOffset);

```

```

15
16     if (sunrise == date && sunset == date) continue; // Polar night
    fallback
17
18     for (DateTime localTime = sunrise; localTime < sunset;
localTime = localTime.AddSeconds(stepSeconds)) {
19         if (!_isSimulating) yield break;
20
21         ApplySunPosition(localTime);
22
23         // Optimization: Skip exact sunrise/sunset steps where
there is functionally no light
24         if (CurrentBeamLux < 0.1f && CurrentDiffuseLux < 0.1f) {
25             continue;
26         }
27
28         // Dispatch hardware raytracing for this specific hour
irradianceBaker.DispatchRays(CurrentSunDirection,
29         CurrentBeamLux, CurrentDiffuseLux, deltaHours, samplesPerPixel,
baker.PositionMap, baker.NormalMap);
30         completedSteps++;
31         currentProgress = ((float)currentDayIdx / totalDays) * 100f
;
32
33         // Extract telemetry data for the CSV asynchronously
TakeAsyncSnapshot(day, localTime.ToString("yyyy-MM-dd HH:mm
34         "), currentAltitude, currentAzimuth, CurrentBeamLux,
CurrentDiffuseLux);
35
36         if (completedSteps % 10 == 0) {
37             ApplyDosePreview();
38         }
39         // Yield execution back to the OS and Unity Main Thread to
maintain interactivity
40         yield return null;
41     }
42
43     // Flush the daily accumulation map to disk as a floating-point
EXR
44     ExportDailyMap(exportDir, day);
45 }
46 }

```

Listing 9.3: LightDoseSimulator.cs: The Coroutine driving dynamic time-stepping, asynchronous telemetry extraction, and synchronous EXR flushing.

Listing 9.4 showcases the background threaded parser utilized to load massively dense geometric heritage assets natively at runtime.

```

1 string path = Path.Combine(ProjectManager.ModelFolder, artifactFile);
2 UpdateProgress(0.1f, $"Loading artifact: {artifactFile}");
3
4 // 1. Offload heavy string parsing to a background CPU thread
5 var parseTask = Task.Run(() => {
6     if (!File.Exists(path)) return null;
7     return SimpleObjLoader.ParseObj(File.ReadAllText(path));
8 });
9
10 // 2. Yield the main thread to keep the "LOADING ASSETS" UI animation

```

```

    smooth
11 while (!parseTask.IsCompleted) {
12     UpdateProgress(0.1f, "Parsing artifact geometry...");
13     yield return null;
14 }
15
16 if (parseTask.IsCompletedSuccessfully && parseTask.Result != null &&
    parseTask.Result.Count > 0) {
17     var dataList = parseTask.Result;
18     _artifactInstance = new GameObject(Path.GetFileNameWithoutExtension
    (artifactFile));
19
20     if (_artifactInstance != null) {
21         _artifactInstance.transform.SetParent(modelRoot);
22         _artifactInstance.transform.localPosition = Vector3.zero;
23         _artifactInstance.transform.localScale = Vector3.one;
24
25         // 3. Assemble the hierarchy on the main thread
26         foreach (var data in dataList) {
27             GameObject child = CreateGameObjectFromData(data,
    artifactMaterial);
28             if (child == null) continue;
29             child.transform.SetParent(_artifactInstance.transform,
    false);
30             child.transform.localPosition = Vector3.zero;
31             SetupArtifactPart(child, project);
32         }
33
34         SetupArtifactRoot(_artifactInstance);
35
36         // 4. Normalize the raw mesh geometry (Centering + Scaling
    children to a 2m box)
37         CenterAndScaleModel(_artifactInstance, 2.0f);
38
39         // 5. Apply user-defined transformations on the root
40         _artifactInstance.transform.localPosition = project.
    artifactPosition;
41         if (project.artifactScale != Vector3.zero)
42             _artifactInstance.transform.localScale = project.
    artifactScale;
43     }
44 }

```

Listing 9.4: RuntimeModelLoader.cs: Multi-threaded ‘.obj’ chunk parsing avoiding main-thread freezes, followed by strict spatial normalization constraints.

9.3. Mathematical Radiometry and Ephemeris Models

Listing 9.5 details the robust celestial mathematics utilized to compute the solar position and air mass attenuation dynamically at runtime.

```

1 // Calculate sun position for a given location and local date/time.
2 public static SunPosition Calculate(double latitude, double longitude,
    double utcOffsetHours, DateTime localTime)
3 {
4     // Convert local time to UTC
5     DateTime utc = localTime.AddHours(-utcOffsetHours);

```



```

6     double jd = ToJulianDay(utc);
7     double T = (jd - 2451545.0) / 36525.0; // Julian centuries
8
9     // Geometric mean longitude & anomaly
10    double L0 = (280.46646 + T * (36000.76983 + T * 0.0003032)) %
360.0;
11    double M = 357.52911 + T * (35999.05029 - 0.0001537 * T);
12    double Mrad = M * Math.PI / 180.0;
13
14    // Equation of center & apparent longitude
15    double C = Math.Sin(Mrad) * (1.914602 - T * (0.004817 + 0.000014 *
T)) + Math.Sin(2 * Mrad) * (0.019993 - 0.000101 * T) + Math.Sin(3 *
Mrad) * 0.000289;
16    double sunLon = L0 + C;
17    double omega = 125.04 - 1934.136 * T;
18    double appLon = sunLon - 0.00569 - 0.00478 * Math.Sin(omega * Math.
PI / 180.0);
19
20    // Solar declination
21    double epsR = (23.439 - 0.013 * T + 0.00256 * Math.Cos(omega * Math
.PI / 180.0)) * Math.PI / 180.0;
22    double decl = Math.Asin(Math.Sin(epsR) * Math.Sin(appLon * Math.PI
/ 180.0));
23
24    // Equation of Time & True Solar Time
25    double e = 0.016708634 - T * (0.000042037 + 0.0000001267 * T);
26    double varY = Math.Tan(epsR / 2.0) * Math.Tan(epsR / 2.0);
27    double eqT = 4.0 * (varY * Math.Sin(2 * L0 * Math.PI / 180.0) - 2 *
e * Math.Sin(Mrad)) * (180.0 / Math.PI);
28    double trueSolar = (utc.TimeOfDay.TotalMinutes + 4.0 * longitude +
eqT) % 1440.0;
29
30    // Hour Angle and Zenith
31    double HArad = (trueSolar / 4.0 - 180.0) * Math.PI / 180.0;
32    double latR = latitude * Math.PI / 180.0;
33    double cosZen = Math.Sin(latR) * Math.Sin(decl) + Math.Cos(latR) *
Math.Cos(decl) * Math.Cos(HArad);
34    double altDeg = 90.0 - (Math.Acos(cosZen) * 180.0 / Math.PI);
35
36    return new SunPosition {
37        AltitudeDeg = (float)altDeg,
38        BeamLux = ClearSkyBeamLux((float)altDeg)
39    };
40 }
41
42 private static float ClearSkyBeamLux(float altitudeDeg) {
43     if (altitudeDeg <= 0f) return 0f;
44     // Kasten-Young formula for optical air mass
45     float airMass = 1f / (Mathf.Sin(altitudeDeg * Mathf.Deg2Rad) +
0.50572f * Mathf.Pow(altitudeDeg + 6.07995f, -1.6364f));
46     return 127500f * Mathf.Pow(0.7f, Mathf.Pow(airMass, 0.678f)); // E
= E0 * 0.7^(AM^0.678)
47 }

```

Listing 9.5: SunCalculator.cs: Calculates solar altitude, azimuth, and direct beam illuminance using the Kasten-Young air mass formula.

Listing 9.6 outlines the Perez Sky model evaluation function utilized by the HLSSL

kernel whenever an indirect ray escapes the geometry bounds.

```

1 // Perez coefficients for clear sky (Turbidity T = 2.0)
2 static const float perez_a = -0.1058; // Horizon brightening/darkening
3 static const float perez_b = -0.0738; // Luminance gradient near
  horizon
4 static const float perez_c = 1.8200; // Circumsolar intensity
5 static const float perez_d = -2.9744; // Circumsolar width
6 static const float perez_e = 0.1778; // Backscattered light
7
8 float PerezFunction(float theta, float gamma)
9 {
10     float cosTheta = max(0.01, cos(theta));
11     float cosGamma = cos(gamma);
12
13     float f = (1.0 + perez_a * exp(perez_b / cosTheta));
14     float g = (1.0 + perez_c * exp(perez_d * gamma) + perez_e *
15         cosGamma * cosGamma);
16
17     return f * g;
18 }
19 float EvaluateSky(float3 direction, float3 sunDir, float beamLux, float
20     diffuseLux, bool includeSun)
21 {
22     if (includeSun && dot(direction, sunDir) > 0.9998) return beamLux;
23
24     // Perez Luminance Distribution
25     float theta = acos(saturate(dot(direction, float3(0, 1, 0))));
26     float gamma = acos(saturate(dot(direction, sunDir)));
27     float relLuminance = PerezFunction(theta, gamma);
28
29     // Normalization factor: Lz ~ Ed * 0.2 for T=2.0
30     return relLuminance * (diffuseLux * 0.2);
31 }

```

Listing 9.6: SkyModel.hlsl: Evaluates the Perez All-Weather Sky Model using 5 empirical coefficients to compute relative luminance for diffuse scattering.

Listing 9.7 demonstrates the robust variance derivation over the telemetry dataset.

```

1 // Extracted from the TakeAsyncSnapshot callback
2 float hourlyVariance = 0f;
3 float expectedArea = (float)activePixels;
4 if (expectedArea > 1f) {
5     // Derived algebraically from Sum(x^2)/N - (Sum(x)/N)^2
6     hourlyVariance = (float)((sumSqrDelta / expectedArea) - (mean *
7         mean));
8     if (hourlyVariance < 0f) hourlyVariance = 0f;
9 }

```

Listing 9.7: LightDoseSimulator.cs: Computes statistical variance across the accumulated dosimetric frame.